

Phantom Pre-Trends and Dynamic Effects in Event-Studies with Imbalanced Panels

Abhinav Ramaswamy*

Nicolas Studen†

Abstract

Two-way fixed-effect event studies are routinely used to test for parallel trends and estimate dynamic treatment effects in difference-in-differences designs. We show that when the panel is unbalanced (i.e. when units either enter or exit the sample in different periods) the event-study estimator fails in two ways. First, under treatment-effect heterogeneity, its pre-trend (dynamic effect) coefficients are biased when units enter (exit) the panel in different periods. The estimator may thus fabricate pre-trends and dynamic effects, or conceal genuine ones. Crucially, this bias exists even when all treated units adopt treatment in the same period — a setting long considered safe, because known problems with the estimator only bite under staggered adoption. Second, because the set of observed units changes across periods, the coefficients compare shifting subpopulations and thus target a misleading estimand. We propose a new estimator that, unlike existing heterogeneity-robust estimators, resolves both problems. We illustrate the empirical relevance of our findings by reanalyzing applied papers in this setting, and providing a new application studying the effect of defections in U.S. House speaker elections on campaign fundraising.

*Ph.D. Candidate, Department of Political Science, Stanford University. acr122@stanford.edu.

†Ph.D. Candidate, Department of Political Science, Stanford University. nstuden@stanford.edu.

1 Introduction

In difference-in-differences (DiD) designs, researchers must gauge the plausibility of the key identification assumption: absent treatment, treated and control units’ average outcomes would have followed parallel paths. Researchers are also often interested in estimating dynamic treatment effects (whether effects grow or shrink over time).

Popularized by landmark economics applications ([Jacobson et al., 1993](#); [Autor, 2003](#)) and recommended by textbooks ([Angrist and Pischke, 2009](#), §5.2), the dynamic event-study regression is the canonical approach across economics and political science for estimating dynamic treatment effects and assessing the parallel trends assumption in panel data. In the simplest setting — a binary treatment with a common onset period t^* — a researcher who observes units $i = 1, \dots, N$ over periods $t = 1, \dots, \mathcal{T}$ estimates:

$$Y_{it} = \alpha_i + \lambda_t + \sum_{k \neq t^* - 1} \beta_k \cdot D_i \cdot \mathbf{1}(t = k) + \varepsilon_{it}, \quad (1)$$

where D_i indicates whether unit i is ever treated, α_i are unit fixed effects, λ_t are period fixed effects, and the period immediately preceding treatment ($k = t^* - 1$) is omitted as the reference. The pre-treatment coefficients $\{\beta_k : k < t^*\}$ are then inspected directly: if they are statistically indistinguishable from zero, the parallel trends assumption is considered justified; if they trend or are jointly significant, parallel trends is called into question. Post-treatment coefficients $\{\beta_k : k > t^*\}$ that differ from the contemporaneous treatment effect β_{t^*} provide evidence of dynamic treatment effects.

What happens, however, when the panel is imbalanced? While most theoretical work on DiD assumes a balanced panel for convenience, many empirical settings feature compositional change: units enter or exit the sample at different times, so the set of observed units varies across periods. Consequently, different units may be observed for a different number of periods. Examples include panels of firms ([Liao, 2023](#); [Latura and Weeks, 2023](#)), of legislators and other officeholders ([Fourinaies and Hall, 2022](#); [Weschle, 2021](#)), and of countries and states

(Acemoglu et al., 2019; Paglayan, 2021), where entry and exit can reflect a unit’s founding or dissolution, an electoral cycle, or the reach of the data source. In each case, the panel is unbalanced, and even when treatment-onset is common, units may be observed for differing numbers of pre- and post-treatment periods.

Applying the canonical event-study estimator (Equation 1), which includes unit fixed effects, to such a panel seems logical. Units entering or exiting at different times likely have different baseline outcome levels. Researchers may want to absorb these differences with unit fixed effects to increase efficiency and address pre-trends or dynamic effects that are driven by the panel’s changing composition rather than within-unit changes.¹

This paper shows that with imbalanced panels, even when treatment onsets in a common period (the simplest case), the canonical estimator is biased under effect-heterogeneity and targets the wrong estimand. The bias stems from how unit fixed effects interact with panel imbalance.

Consider a setting with common treatment-onset and where pre- and post-treatment outcomes are constant within-units (i.e. there are no parallel trends violations or dynamic effects). Unit fixed effects demean each unit’s outcome with respect to its *entire* observed history, including post-treatment periods. A treated unit’s demeaned outcome is therefore contaminated by its treatment effect. The magnitude of this contamination depends on both the unit’s effect-size and the share of its observations that fall after treatment. Because this share varies across entrance- or exit-cohorts (units which enter or exit in a given period), treated units’ demeaned outcomes differ even when they have identical raw outcomes. As the mix of observed cohorts changes from period to period, the average shifted-outcome drifts across periods, even when average raw outcomes do not. Under differential-entry, the drift appears in the pre-treatment periods, and under differential-exit, it appears in the post-treatment periods.

¹Because the composition of observed units changes from period to period, level differences between cohorts can themselves generate movement in the pooled outcome means over the pre-treatment periods (in settings with differential-entry) or post-treatment periods (in settings with differential exit), even without any within-unit trends.

When average effects are heterogeneous across cohorts, these mechanical drifts result in “phantom” pre-trend or dynamic effect coefficients. A phantom pre-trend is a non-zero pre-treatment coefficient that appears even when parallel trends hold exactly, while a phantom dynamic effect is a post-treatment coefficient that grows or shrinks relative to the treatment effect estimate, even when no such dynamic effects exist. This distortion can also run in the opposite direction, masking pre-trends and dynamic effects that genuinely exist. For applied researchers, this means that a failed pre-trend diagnostic in an unbalanced panel need not condemn a design. The failure can lie in the estimator, or in the population the coefficients describe, not in the assumption itself. Importantly, the coefficients are unbiased when effects are homogeneous across cohorts. We show these results extend to the staggered treatment-adoption setting as well.

Independent of this bias, the event-study interactions target the wrong estimand for assessing pre-trends and estimating dynamic effects. Under differential-entry,² pre-treatment interactions for a period before all cohorts have entered are estimated only from the cohorts present in that period, not the full set of units used to estimate the treatment effect. They therefore do not test parallel trends for all the units that ultimately identify the treatment effect. Symmetrically, under differential-exit,³ a post-treatment coefficient for a period after some cohorts have exited is identified only from the survivors — a different subset than the contemporaneous coefficient β_{t^*} , which targets all treated units. Later post-treatment coefficients are thus not directly comparable to β_{t^*} , and a researcher may mistake the changing composition of surviving cohorts for growing or shrinking treatment effects relative to β_{t^*} . [Liu \(2025\)](#) shows an analogous version of this estimand issue exists in event-studies with staggered adoption and panel balance — aggregated “lead and lag” coefficients are not directly comparable, because they are identified from changing subsets of units.

Recent “heterogeneity-robust” (hereafter: HTE-robust) estimators ([Callaway and Sant’Anna, 2021](#); [Sun and Abraham, 2021](#); [Borusyak et al., 2024](#); [Liu et al., 2024](#); [Baker et al., 2022](#))

²i.e. when units enter the data in different pre-treatment periods

³i.e. when units exit the data in different post-treatment periods

do not fully resolve the issues we identify. These estimators were developed to handle the “forbidden-comparisons” bias,⁴ which occurs *only* under staggered adoption. Crucially, this is not the mechanism that drives the bias result in this paper, as evidenced by the fact that our bias extends to the non-staggered treatment-adoption setting. Because these estimators do not directly model compositional change, the estimand issue remains, and we show that in either treatment-timing regime, they may generate spurious pre-trend or dynamic effect coefficients.

Existing approaches cannot resolve both the bias and estimand issues in imbalanced panels. To tackle this, we redefine the estimand as a weighted average of the *within-cohort* treatment effects, where a cohort is the set of units that enter and exit the panel in the same periods. Because a cohort has fixed composition, this restores, within each cohort, a well-defined pre-trend test and dynamic-effect target when the overall sample exhibits differential entry or exit, respectively. It also implies a corresponding estimation approach that is unblemished by the phantom bias: estimating the main effect, dynamic effects, and diagnosing parallel trends, within cohorts.

To achieve parallel trends within cohorts, we recommend matching treated to control units in a cohort based on pre-treatment outcome trajectories and covariates, and then estimating an event-study without unit fixed effects (our approach also accommodates inverse-propensity weighting and outcome modeling methods).⁵ Beyond resolving the bias and estimand problems, this strategy inherits the usual benefits of matching: it addresses *genuine* pre-trend violations driven by compositional changes or within-unit trends, makes parallel trends in the treatment period more credible by balancing covariates across treated and con-

⁴Under staggered adoption, TWFE estimators identify effects for each treatment-cohort using already-treated and later-treated units as controls (“forbidden comparisons”) (Goodman-Bacon, 2021; Sun and Abraham, 2021), which allows the weights that aggregate these effects into a point estimate to be negative (often referred to as the “negative-weighting problem”). Consequently, when effects are heterogeneous across units, it is possible that every unit has a positive treatment effect in each period but the model still reports an overall negative effect (de Chaisemartin and D’Haultfœuille, 2020).

⁵Because each cohort is a balanced sub-panel, one could estimate the canonical event-study within-cohort. However, some sort of matching will likely be necessary to achieve parallel trends, after which, including unit fixed effects is inefficient.

trol units, and reduces residual variance. Furthermore, for pre-trends, we recommend *against* aggregation: parallel trends is clearest assessed cohort by cohort. For dynamic effects, we recommend aggregating the cohort-specific coefficients at each event time, weighting each cohort by its share of the treated units observed in that period — the same weighting as the main effect. Once differential exit sets in, we apply the same rule but report the share of treated units that have left the sample, so the plot is explicit that the estimand localizes to the surviving cohorts from that period onward.⁶

We illustrate the practical stakes of our findings and proposed estimator in two ways. First, we document several published DiD designs that share our setting of an unbalanced panel with differential entry (Liao, 2023; Weschle, 2021; Clarke, 2020; Fourinaies and Hall, 2022; Latura and Weeks, 2023). We reanalyze one in the main text and two more in the supplemental appendix, and show that in each case a more precisely defined estimand would support a more convincing identification story. In the main-text example, we show how a properly targeted estimation approach changes a published substantive result (Liao, 2023). Second, we apply the estimator to a new study of the effect of defections in U.S. House speaker elections on campaign fundraising, where the pooled diagnostic would have led a careful researcher to abandon a design that within-cohort analysis shows to be sound.

2 The Bias of the Canonical Event-Study Estimator

In this section, we prove the bias of the canonical event-study estimator (hereafter, ES-estimator, see Equation 1) for imbalanced panels. Specifically, we show that when treatment effects are heterogeneous across cohorts, the ES-estimator can generate the appearance of pre-trend violations even when parallel trends holds exactly. The opposite concern will also hold: the ES-estimator may report no pre-trend violations when parallel trends breaks.

⁶This is not equivalent to running an off-the-shelf event study and annotating it with attrition counts: estimating within cohorts first is what keeps composition change from biasing the dynamics. For example, were cohorts to sit at different post-treatment baselines and exit one by one, with no unit actually trending, a pooled estimator would manufacture growing dynamic effects, whereas within-cohort aggregation would not.

2.1 Setup

We work throughout in the *single-treatment-onset* (non-staggered) setting: every treated unit receives treatment at a common period t^* , and the event study takes the form of Equation (1). We develop the result in the *differential-entry* case — units enter the panel in different pre-treatment periods — as this setting threatens identification of the main effect, which is of greatest interest to practitioners. For clarity, we work with *common-exit* — units are observed through a common terminal period. However, analogous results apply to settings with *differential-exit*, as the mechanism will depend only on units being observed for differing numbers of periods.

Notation. We observe units $i = 1, \dots, N$ over periods $t = 1, \dots, \mathcal{T}$, with common treatment onset at t^* . Each unit is either treated ($D_i = 1$) or a control ($D_i = 0$). We partition units into C *entry cohorts*, indexed by $c \in \{1, \dots, C\}$ and defined by the period e_c in which the cohort enters the panel, with

$$e_1 < e_2 < \dots < e_C < t^*.$$

A cohort thus collects all units — treated and control — that first appear in the panel in the same period.⁷ Because all units are observed through the common terminal period $\mathcal{T}$, cohort c is observed for

$$T_{\text{pre}}(c) = t^* - e_c \quad \text{pre-treatment periods}, \quad T_{\text{post}} = \mathcal{T} - t^* + 1 \quad \text{post-treatment periods},$$

for a total of $T_c = T_{\text{pre}}(c) + T_{\text{post}}$ observed periods. Later-entering cohorts have smaller $T_{\text{pre}}(c)$ and hence smaller T_c (this is the panel imbalance), while T_{post} is common across cohorts.

⁷Both treated and control units belong to cohorts; what defines a cohort is entry timing, not treatment status. Each cohort contains n_c treated and m_c control units.

Two notions of parallel trends. Because the composition of observed units changes across periods, “parallel trends” admits two distinct meanings that coincide in a balanced panel but must be separated here:

(PT-W) **Within-cohort parallel trends.** For every cohort c , treated and control units in that cohort share a common counterfactual trend: $\mathbb{E}[Y_{it}(0) - Y_{i,t-1}(0) \mid D_i = 1, c] = \mathbb{E}[Y_{it}(0) - Y_{i,t-1}(0) \mid D_i = 0, c]$ for all $t \in \{e_c + 1, \dots, t^*\}$. This is the identification condition relevant to the treatment effect each cohort contributes.

(PT-P) **Pooled parallel trends.** Averaging over all units observed in each period, the mean treated–control outcome difference is constant across pre-treatment periods: $\mathbb{E}[Y_{it}(0) - Y_{i,t-1}(0) \mid D_i = 1] = \mathbb{E}[Y_{it}(0) - Y_{i,t-1}(0) \mid D_i = 0]$ for all $t \in \{1, \dots, t^*\}$.

Data-generating process. To prove the bias, we study a generic DGP in which both notions of parallel trends hold exactly by construction:

$$Y_{it} = \begin{cases} \mu_0 & \text{if } D_i = 0, \\ \mu_1 & \text{if } D_i = 1 \text{ and } t < t^*, \\ \mu_1 + \tau_c & \text{if } D_i = 1 \text{ and } t \geq t^*, \end{cases} \quad (2)$$

where μ_1 and μ_0 are scalar constants so there is a fixed level difference between treated and control units, and τ_c is a *cohort-specific* treatment effect. Outcomes are flat in the pre-period (allowing for common time trends does not change the result, see Appendix A), and the treatment effect is static (no dynamics). Thus, the only departure from a textbook two-group design is that the panel is unbalanced. Under this DGP both PT-W and PT-P hold exactly: within any cohort the treated–control gap is $\mu_1 - \mu_0$ in every pre-period, and the pooled treated and control pre-period means are flat at μ_1 and μ_0 regardless of which cohorts happen to be observed. The result below is therefore striking: the canonical estimator manufactures pre-trends even though there is no pre-trend to detect under either

notion.

2.2 How the Bias Emerges

By the Frisch-Waugh-Lovell (FWL) theorem, the event-study coefficients $\hat{\beta}_k$ in Equation (1) are identical to those from regressing the *unit-demeaned* outcome $\tilde{Y}_{it} \equiv Y_{it} - \bar{Y}_i$ on the unit-demeaned interaction regressors and demeaned time dummies, where $\bar{Y}_i$ averages over *all* of unit i 's observed periods. The key observation is that this average runs over the unit's *entire* history — including post-treatment periods.

Demeaned outcome. For a treated unit in cohort c , the time-average outcome is

$$\bar{Y}_i = \frac{T_{\text{pre}}(c) \mu_1 + T_{\text{post}} (\mu_1 + \tau_c)}{T_c} = \mu_1 + \frac{T_{\text{post}}}{T_c} \tau_c, \quad (3)$$

so the treatment effect contaminates the mean. The demeaned outcome is therefore nonzero in every period:

$$\tilde{Y}_{it} = \begin{cases} -\frac{T_{\text{post}}}{T_c} \tau_c, & t < t^*, \\ \frac{T_{\text{pre}}(c)}{T_c} \tau_c, & t \geq t^*, \end{cases} \quad (4)$$

with $\tilde{Y}_{it} = 0$ for all control units. Crucially, the pre-period value depends on T_c , so it *varies across cohorts*. As later pre-periods admit later-entering cohorts (larger T_{post}/T_c), the cross-sectional average of $\tilde{Y}$ shifts from period to period — a “trend” on the left-hand side that is purely an artifact of demeaning.

Demeaned regressors. Consider a post-treatment interaction $X_{it}^{(k)} = D_i \cdot \mathbf{1}(t = k)$ for $k \geq t^*$. For a treated unit in cohort c (observed in every post-period, since $e_c < t^* \leq k$), demeaning gives $\tilde{X}_{it}^{(k)} = 1 - 1/T_c$ at $t = k$ and $-1/T_c$ otherwise. Summing over all post-

treatment interactions,

$$\sum_{k \geq t^*} \tilde{X}_{it}^{(k)} = \begin{cases} -\frac{T_{\text{post}}}{T_c}, & t < t^*, \\ \frac{T_{\text{pre}}(c)}{T_c}, & t \geq t^*, \end{cases} \quad (5)$$

which mirrors $\tilde{Y}_{it}$ in (4) period by period, governed by the *same* factor T_{post}/T_c .

Homogeneous effects: no bias. If $\tau_c = \tau$ for all c , then comparing (4) and (5) gives

$$\tilde{Y}_{it} = \tau \cdot \sum_{k \geq t^*} \tilde{X}_{it}^{(k)} \quad \text{for every observed } (i, t). \quad (6)$$

The demeaning distortion on the left is absorbed exactly by the post-treatment regressors on the right with the common coefficient τ . Nothing is left over for the pre-treatment interactions.

Heterogeneous effects: bias emerges. If τ_c varies across cohorts, the ratio $\tilde{Y}_{it} / \sum_{k \geq t^*} \tilde{X}_{it}^{(k)} = \tau_c$ is cohort-specific, but OLS constrains the post-treatment coefficients to be common, so the post-treatment interactions cannot reproduce $\tilde{Y}$ for all cohorts at once. Fitting the post-interactions with a common value $\bar{\tau}$ therefore leaves, for a treated unit in cohort c in a pre-period, the residual

$$\tilde{Y}_{it} - \bar{\tau} \sum_{k \geq t^*} \tilde{X}_{it}^{(k)} = -\frac{T_{\text{post}}}{T_c} \delta_c, \quad \delta_c \equiv \tau_c - \bar{\tau}, \quad (7)$$

which is constant within a cohort but differs across cohorts. Averaging this residual over the treated units observed in pre-period k — the cohorts that have entered by then, $\{c : e_c \leq k\}$ — gives

$$\bar{r}_k = -\frac{\sum_{c: e_c \leq k} n_c \frac{T_{\text{post}}}{T_c} \delta_c}{\sum_{c: e_c \leq k} n_c}. \quad (8)$$

As k runs from the earliest entry e_1 toward $t^* - 1$, the set $\{c : e_c \leq k\}$ expands as later cohorts enter, so $\bar{r}_k$ generically changes from one pre-period to the next. This period-

to-period movement is the phantom: among treated units, the pre-treatment interactions $D_i \cdot \mathbf{1}(t = k)$ are the only regressors that distinguish one pre-period from another, so they must absorb the shifting residual, yielding $\hat{\beta}_k \neq 0$. Only when effects are homogeneous ($\delta_c = 0$ for all c) is $\bar{r}_k$ constant at zero and the phantom absent, recovering the cancellation of the previous paragraph.

2.3 Formal Results

The mechanics above imply that in imbalanced panels with differential-entry, phantom pre-trends emerge under effect heterogeneity across cohorts.

Proposition 1 (No phantom pre-trends under homogeneous effects). *When $\beta_k = 0$ for every pre-treatment period $k < t^*$ (parallel trends holds exactly both in the pooled (PT-P) and within-cohort (PT-W) senses) and $\tau_c = \tau$ for all c (treatment effects are homogeneous across cohorts c), then $\hat{\beta}_k = \beta_k = 0$ for each $k < t^*$, regardless of the panel composition (the cohort entry times e_c and sizes n_c, m_c).*

Proof. By (4) and (5), $\tilde{Y}_{it} = \tau \sum_{k \geq t^*} \tilde{X}_{it}^{(k)}$ for every observed (i, t) (for treated units by direct computation; for controls both sides are zero). Hence the coefficient vector $\beta_k = \tau$ for $k \geq t^*$, $\beta_k = 0$ for $k < t^*$, and $\lambda_t = 0$ produces a perfect fit (residual sum of squares equal to zero). As this is the global minimum and the OLS solution is unique under full column rank, $\hat{\beta}_k = 0$ for all $k < t^*$. $\square$

Proposition 2 (Phantom pre-trends under heterogeneous effects). *When $\beta_k = 0$ every pre-treatment period $k < t^*$ (parallel trends holds exactly in both the pooled (PT-P) and within-cohort (PT-W) senses) and $\tau_c \neq \tau_{c'}$ for some $c \neq c'$ (treatment effects are heterogeneous across cohorts c), then $\exists k < t^*$ such that $\hat{\beta}_k \neq \beta_k = 0$.*

Proof Sketch. See Section 2.2. In short, it shows that the average residual among treated units in the pre-periods ($\bar{r}_k$ in (8)) varies across pre-periods whenever effects are heterogeneous across cohorts. Among treated units the pre-treatment interactions are the only

regressors that distinguish one pre-period from another, so they must absorb this period-varying residual, giving $\hat{\beta}_k \neq 0$ for some $k < t^*$.

Formal Proof. See Appendix A which derives the exact coefficients $\hat{\beta}_k = \bar{a} - \bar{a}_{S_k}$ (which is also the bias, since $\beta_k = 0$), where $\bar{a}$ is the average treated cohort effect and $\bar{a}_{S_k}$ its average over the cohorts observed at period k .

Corollary 1 (Obscuring genuine pre-trends). *Proposition 2 implies that the event-study estimator can also obscure genuine pre-trends.*

Proof. Decompose the outcome into its untreated potential outcome and the realized effect,

$$Y_{it} = Y_{it}(0) + D_i \mathbf{1}(t \geq t^*) [Y_{it}(1) - Y_{it}(0)].$$

Let $\text{ES}_k(w)$ denote the k th-period coefficient that the canonical event-study estimator returns when run on outcome w . Since OLS is linear in the outcome,

$$\hat{\beta}_k = \text{ES}_k(Y) = \underbrace{\text{ES}_k(Y(0))}_{\equiv \beta_k} + \underbrace{\text{ES}_k(D_i \mathbf{1}(t \geq t^*) (Y_{it}(1) - Y_{it}(0)))}_{\equiv b_k},$$

where β_k is the genuine pre-trend the test aims to detect. In our setup $Y_{it}(1) - Y_{it}(0) = \tau_{c(i)}$, so b_k is the event study applied to the outcome $W_{it} \equiv D_i \mathbf{1}(t \geq t^*) \tau_{c(i)}$, which is zero in every pre-period and for every control unit, and equals the treatment effect $\tau_{c(i)}$ only for treated units in post-treatment periods. By Proposition 2, this produces nonzero pre-treatment coefficients ($b_k \neq 0$ for some $k < t^*$) whenever effects are heterogeneous across cohorts c . Because b_k depends only on the treatment effects and the panel structure — not on $Y_{it}(0)$ — it is *added* to whatever pre-trend is present. When $\beta_k = 0$ it manufactures a violation ($\hat{\beta}_k = b_k \neq 0$); when $\beta_k \neq 0$ it can inflate, attenuate, or — if $b_k \approx -\beta_k$ — cancel it, so the estimator can report a null precisely when parallel trends fails. $\square$

The canonical estimator therefore does not merely add spurious pre-trends; it is biased for the pre-trend by the fixed amount b_k , in either direction. We derive this exact bias in

Appendix A. Crucially, we take the pooled target PT-P ($\beta_k = \text{ES}_k(Y(0))$) for concreteness, but composition change makes even this pooled coefficient an imperfect summary of the underlying trends; we return to a cohort-specific target when we introduce our estimator.

3 Simulation Evidence

This section provides numerical illustrations of the phantom biases across different data generating processes.

3.1 Phantom Pre-Trends in the Proof’s Data Generating Process

We simulate the DGP of Equation (2) over a panel of six periods with treatment onset in the fourth (event time $k = 0$), so the omitted reference is event time -1 . Three entry cohorts enter in the first three periods and are observed through the end of the panel; they are thus observed for $T_{\text{pre}} = 3, 2, 1$ pre-treatment periods, respectively, with a common $T_{\text{post}} = 3$. Each cohort contains 30 treated and 30 control units. We set the control mean to $\mu_0 = 1$ and the treated pre-treatment mean to $\mu_1 = 2$. To give unit fixed effects a genuine efficiency rationale, we add a persistent unit-level intercept $\alpha_i \sim N(0, \sigma_\alpha^2)$ with $\sigma_\alpha = 1$ (constant across each unit’s periods) and small idiosyncratic noise $\varepsilon_{it} \sim N(0, 0.1^2)$; both within-cohort (PT-W) and pooled (PT-P) parallel trends continue to hold in expectation, so there is no pre-trend in the data. We fit the canonical event study of Equation (1) and plot the estimated coefficients $\hat{\beta}_k$ against event time.

Figure 1 shows the homogeneous treatment effect case, $\tau_c = -6$ for all three cohorts. The pre-treatment coefficients are zero despite the unbalanced panel, confirming Proposition 1: panel imbalance alone does not generate a phantom pre-trend.

Figure 2 introduces heterogeneous effects: $\tau = (-10, -6, -2)$ for the cohorts observed for $T_{\text{pre}} = 3, 2, 1$ pre-periods, respectively. The same estimator now produces large, statistically significant pre-treatment coefficients that decline monotonically toward the reference period




Figure 1: Homogeneous treatment effects: no phantom pre-trend.

Notes: Event-study coefficients $\hat{\beta}_k$ from the canonical specification of Equation (1). Cohort effects are homogeneous at $\tau_c = -6$. The DGP satisfies parallel trends exactly, so the true value of every pre-treatment coefficient is zero. Bars are 95% confidence intervals. The dotted line marks treatment onset, and event time -1 is the omitted reference.

($\approx 2.7 \rightarrow 1.5 \rightarrow 0$) — a downward phantom pre-trend — even though the data contain no pre-trend whatsoever, confirming Proposition 2. This pattern encourages the researcher to wrongly reject parallel trends or to dismiss a real effect as spurious.

3.2 Alternative DGPs

The phantom is not an artifact of the specific data-generating process used in the proof. The subsections below each modify that DGP along one dimension — adding a common time trend, moving the imbalance to the post-period, and letting cohorts differ in baseline level — and show the distortion persists. Throughout, we keep the within-cohort unit-level variance of Section 3.1, so unit fixed effects always carry an efficiency rationale.




Figure 2: Heterogeneous treatment effects: a downward phantom pre-trend.

Notes: As Figure 1, with heterogeneous cohort effects $\tau = (-10, -6, -2)$. The earliest-entering cohort has the largest-magnitude effect. The DGP again satisfies parallel trends exactly, so the true value of every pre-treatment coefficient is zero. The estimated coefficients decline toward the reference period and are individually distinguishable from zero. Bars are 95% confidence intervals. The dotted line marks treatment onset, and event time -1 is the omitted reference.

3.2.1 Common Time Trend

The flat DGP of Equation (2) held outcomes constant in the pre-period. We now relax this, adding a common time effect γ_t shared by treated and control units:

$$Y_{it} = \begin{cases} \mu_0 + \gamma_t & \text{if } D_i = 0, \\ \mu_1 + \gamma_t & \text{if } D_i = 1 \text{ and } t < t^*, \\ \mu_1 + \gamma_t + \tau_c & \text{if } D_i = 1 \text{ and } t \geq t^*. \end{cases} \quad (9)$$

The time effect γ_t may take any shape (e.g. linear or non-monotone) and is common to both groups, so the treated–control difference remains $\mu_1 - \mu_0$ in every pre-period: both within-cohort (PT-W) and pooled (PT-P) parallel trends continue to hold exactly. We do not allow γ_t to vary across cohorts, as cohort-specific time effects would break PT-P and so confound the exercise; here outcomes simply move over time, with no parallel-trends violation of any kind.

We simulate (9) with an arbitrary non-monotone γ_t , keeping all other features of Section 3.1 fixed — the same three cohorts, the same heterogeneous effects $\tau = (-10, -6, -2)$, and the same unit-level variance. Figure 3 reports the result. Panel (a) confirms the trend is genuinely present in the raw data: both groups move sharply over time, yet track each other in parallel through the pre-period (a constant gap of $\mu_1 - \mu_0 = 1$). Panel (b) shows that the event study — whose period fixed effects λ_t absorb γ_t exactly — produces pre-treatment coefficients identical to the flat-DGP case ($\approx 2.7 \rightarrow 1.5 \rightarrow 0$). The phantom downward pre-trend is entirely unaffected by the time trend, as established analytically in Appendix A.

3.2.2 Same-Entry, Differential-Exit

The mechanism depends on variation in panel *length*, not on entry timing specifically, and it is symmetric across the two ends of the panel: the estimator’s coefficients are distorted wherever the composition of observed cohorts changes. To show this, we move the imbalance entirely




Figure 3: Arbitrary common time trend: the phantom pre-trend is unchanged.

Notes: DGP (9) with a non-monotone common time effect γ_t shared by treated and control units. Panel (a) plots raw group means by event time and panel (b) event-study coefficients $\hat{\beta}_k$, from the same specification as Figure 2. The period fixed effects absorb γ_t , so the pre-treatment coefficients are those of Figure 2. The dotted line marks treatment onset, and event time -1 is the omitted reference.

to the post-period. All units now enter at a common period — so every unit is observed for the same number of pre-treatment periods T_{pre} — but they *exit* at different times. A cohort c is now defined by its exit period and observed for a cohort-specific number of post-treatment periods $T_{\text{post}}(c)$, with $T_c = T_{\text{pre}} + T_{\text{post}}(c)$. The outcome process is otherwise Equation (2), with static, cohort-specific effects τ_c , so both PT-W and PT-P again hold exactly and *no cohort's effect changes over time*.

We simulate three exit cohorts observed for $T_{\text{post}}(c) = 1, 2, 3$ post-treatment periods (common $T_{\text{pre}} = 3$) with heterogeneous effects $\tau = (-2, -6, -10)$, keeping the same unit-level variance as before. Figure 4 shows the result. The pre-treatment coefficients are now flat — with a common entry period, the pre-period composition never changes — but the *post-treatment* coefficients trend steeply (from ≈ -6 at onset to ≈ -8.7 two periods later), even though every cohort's effect is static. The composition change has simply moved to the post-period, where cohorts drop out one by one, and the estimator manufactures dynamic effects exactly as it manufactured pre-trends under differential entry.




Figure 4: Same entry, differential exit: a phantom in the dynamic effects.

Notes: Event-study coefficients $\hat{\beta}_k$ from the canonical unit-fixed-effects specification. All units enter at a common period (common $T_{\text{pre}} = 3$) but exit after $T_{\text{post}}(c) = 1, 2, 3$ post-treatment periods, with static heterogeneous effects $\tau = (-2, -6, -10)$; both PT-W and PT-P hold. The pre-treatment coefficients are flat, but the post-treatment coefficients trend despite no cohort's effect changing over time, because the composition of observed cohorts changes across the post-period. Bars are 95% confidence intervals; the dotted line marks treatment onset; event time -1 is the omitted reference.

3.2.3 Baseline Differences Across Cohorts

In the settings so far, both notions of parallel trends held. We now let cohorts sit at *different baseline levels*, so that pooled parallel trends (PT-P) fails while within-cohort parallel trends (PT-W) still holds and no unit’s outcome changes in the pre-period. We return to the differential-entry design of Section 3.1 but give each entry cohort a different treated–control gap, holding the heterogeneous effects $\tau = (-10, -6, -2)$ fixed. Because the composition of observed cohorts changes across the pre-period, the pooled treated–control difference drifts even though each cohort’s gap is constant (Figure 5(a)): PT-P is violated by composition alone.

This is the setting in which unit fixed effects appear most defensible — a researcher would include them not only for efficiency but to absorb the cross-cohort baseline differences that make the pooled comparison non-parallel. Yet the unit-fixed-effects event study still returns a phantom pre-trend (Figure 5(b)), of essentially the same magnitude as the flat-baseline case ($\approx 2.7 \rightarrow 1.5 \rightarrow 0$). The fixed effects difference the baselines out, so the reported pre-trend is neither the pooled PT-P violation (which they remove) nor a within-unit trend (there is none): it is the effect-heterogeneity phantom, invariant to the baselines (Appendix A).

4 The “Estimand” Problem

Event absent the bias under panel imbalance and heterogeneous effects, the event-study estimator’s coefficients are not directly comparable: they target an estimand that is uninformative about the identification assumption.

Consider the common treatment-onset time and differential-entry setting. A pre-treatment coefficient for a period before all cohorts have entered compares the units observed in that period against the larger set of units observed in the reference period. It therefore does not describe how the *full* set of treated and control units that identify the ATT were trending relative to one another before onset, and therefore does not provide evidence of whether




Figure 5: Baseline differences across cohorts: PT-P fails, yet the phantom is unchanged.

Notes: The differential-entry design of Figure 2 with cohort-specific baseline levels (a treated–control gap that varies across entry cohorts) and the same heterogeneous effects $\tau = (-10, -6, -2)$. Panel (a): raw pooled group means. The pooled treated–control gap changes across the pre-period purely because the composition of observed cohorts changes — pooled parallel trends (PT-P) fails — while within each cohort the gap is constant and no unit trends (PT-W holds). Panel (b): event-study coefficients from the unit-fixed-effects specification; the pre-treatment coefficients are the same phantom as Figure 2, because the unit fixed effects difference the baselines out. The dotted line marks treatment onset; event time -1 is the omitted reference.

those treated units would, absent treatment, have tracked the controls on average. What the coefficient targets instead is pooled parallel trends (PT-P) — the treated–control mean difference among whoever is observed, period by period — which under compositional change is neither necessary nor sufficient for identification and can be actively misleading. A zero pre-treatment coefficient, for example, may mean that treated and control units genuinely track one another, or it may merely reflect composition change: treated and control units’ paths are diverging, but as new units enter, the difference in means among those present in each period stays constant relative to the reference. The coefficient alone cannot distinguish the two. Figure 6 provides an example visualization of this.

The analogous problem applies to the dynamic-effect coefficients under differential exit. A post-treatment coefficient for a period after some cohorts have exited is identified only from the survivors — a smaller, different set of units than the contemporaneous effect $\hat{\beta}_{t^*}$, which is identified from all treated units. Successive lags therefore describe different populations,




Figure 6: Entering cohorts can mask a genuine pre-trend, leaving a zero coefficient.

Notes: A stylized differential-entry design in which controls are flat and cohorts are evenly sized, so each line is one cohort's treated-control mean difference (labeled by entry period). Only the cohort observed throughout (entry -4) is trending — a genuine differential pre-trend — while each later-entering cohort is flat and enters at a progressively lower gap. Because those entrants pull the cross-cohort average down by exactly the right amount, the pooled difference that the event study targets (black) is flat across the entire pre-period, so every pre-treatment coefficient is zero. The entering cohorts mask the one cohort's genuine trend. Event time -1 is the reference period.

and a researcher may read the changing composition of surviving cohorts as an effect that grows or shrinks over time when no unit’s effect is in fact changing (Section 3.2.2).

This comparability problem is not unique to imbalanced panels. Liu (2025) shows an analogous issue in the leads-and-lags event study under *staggered* adoption with a *balanced* panel: because treatment cohorts are observed over different numbers of pre- and post-treatment periods in event time, the aggregated lead and lag coefficients are estimated from changing subsets of cohorts and are likewise not directly comparable. We show that a related problem arises from compositional change even in the non-staggered adoption setting typically regarded as safe.

5 How do Existing HTE-Robust Estimators Perform?

The HTE-robust estimators were built for staggered adoption, so it is natural to ask whether the machinery that makes them robust to cross-cohort effect heterogeneity also protects them against the imbalance problems of this paper. It does not. We show by simulation that when adoption is staggered *and* each adoption cohort is internally imbalanced, both failures reappear—and, revealingly, each estimator fails for reasons unrelated to the forbidden-comparison problem it was designed to solve.

Design. We simulate a staggered panel in which treatment turns on at two dates, $g \in \{g_1, g_2\}$, alongside a never-treated control group. We refer to the set of units first treated at a common date g as an *adoption cohort*, to distinguish it from the *entry cohorts* of Section 2.2, which are defined by when a unit enters the *panel*. The crucial feature is that each adoption cohort is itself an unbalanced panel of exactly the form studied above: within an adoption cohort, units belong to several entry cohorts (differential entry, common exit) or several exit cohorts (common entry, differential exit), and effects may be heterogeneous across those sub-cohorts. Parallel trends holds throughout in both the within-cohort (PT-W) and pooled (PT-P) senses, treatment effects are static, and outcomes are flat in the pre-period, so any

non-zero pre-trend or post-treatment dynamic in the estimated coefficients is spurious. We add modest idiosyncratic noise and use large cohorts, so the reported coefficients coincide with their estimands to two decimals; the results are algebraic properties of the design, not sampling artifacts.⁸

The organizing observation is that all four estimators partition the data by *adoption timing*, not by *entry or exit timing*. Within a single adoption cohort, the differential-entry/heterogeneous-effect structure of Section 2.2 is reproduced exactly, and no amount of cross-adoption-cohort aggregation can undo a bias that lives inside each adoption cohort’s own event study. The simulations below confirm this and separate two distinct channels of failure: the *demeaning phantom* of Propositions 1–2, and the *composition* (estimand) problem of Section 4.

Sun–Abraham and stacked DiD reproduce the demeaning bias. The interaction-weighted estimator (Sun and Abraham, 2021) and stacked DiD (Baker et al., 2022) both estimate each adoption cohort’s event-study path with a unit-fixed-effects (equivalently, unit-demeaned) regression. Because that adoption cohort is internally a differential-entry panel with heterogeneous τ_c , the demeaning phantom of Proposition 2 arises inside each cohort-specific path and survives aggregation. Panel A of Table 1 shows the result: both estimators report a pronounced downward pre-trend (coefficients running from -2.3 toward zero as the reference period approaches) even though no unit trends and PT holds exactly. The dual under differential exit is a phantom *dynamic* effect: Panel A of Table 2 shows post-treatment coefficients that drift with the event time although every unit’s effect is static. A falsification check confirms the mechanism—eliminating either the within-cohort entry heterogeneity or the cross-cohort effect heterogeneity returns all pre-period coefficients to zero. These estimators were designed to purge forbidden comparisons across adoption cohorts; they offer no protection against imbalance *within* one.

⁸Full replication code, including a falsification battery, is in the replication package.

Callaway–Sant’Anna and imputation avoid the demeaning bias. Neither the Callaway–Sant’Anna estimator (Callaway and Sant’Anna, 2021) nor the imputation estimator (Borusyak et al., 2024; Liu et al., 2024) demeans over post-treatment outcomes: the former forms base-period differences against clean controls, the latter fits the untreated potential outcome on control observations only and imputes the counterfactual. The contaminated baseline of Equation (4) therefore never enters, and both are clean in the differential-entry, heterogeneous-effect design that breaks Sun–Abraham (columns 3–4, Panel A of Table 1: pre-trend coefficients are zero, the static effect is recovered exactly). This is the sense in which they “resolve part of the bias”—but, as their construction targets forbidden comparisons rather than the cohort structure, the improvement is incidental to what they were built for.

The estimand problem persists for all four estimators. Avoiding the phantom does not make the coefficients comparable across event time. Under differential exit with heterogeneous effects, both estimators trace the survivor-average effect, which drifts mechanically as short-lived cohorts exit and longer-lived ones remain (Panel A of Table 2, columns 3–4: the aggregate falls from -6 to -10 although no unit’s effect changes). Each coefficient is the correct effect for the units present at that horizon; the drift is the estimand problem of Section 4, not a phantom.

Callaway–Sant’Anna is also exposed to a second, sharper form of the composition problem. When cohorts differ in baseline *levels*—so that PT-W holds but PT-P fails—the estimator reports spurious pre-trends under differential entry and spurious dynamic effects under differential exit, *even with homogeneous effects and no unit trend whatsoever*. Panel B of Tables 1 and 2 shows these coefficients reproducing the naive, no-unit-fixed-effects pooled event study almost exactly. The reason is that with an unbalanced panel the estimator compares changing-composition group *means* across periods rather than within-unit changes: it correctly measures the treated–control differential among the units observed in each period,

but as that set of units turns over, a fixed level gap across sub-cohorts is read off as a trend. The imputation estimator is immune here, because its counterfactual model carries unit fixed effects and so differences out baseline levels; it returns zero pre-trends and a flat effect throughout (final columns of Panels B).

Summary. Table 3 collects the pattern. The two biases are governed by orthogonal triggers. The *demeaning phantom* requires effect heterogeneity interacted with imbalance and strikes the unit-demeaning estimators (Sun–Abraham, stacked DiD); the *composition* bias requires baseline-level differences interacted with imbalance and strikes the estimator that compares group means across a shifting sample (Callaway–Sant’Anna). The imputation estimator is the most robust of the four—immune to both biases—and its only residual weakness is the survivor-conditional estimand under differential exit, which is unavoidable for any estimator that reports a single dynamic path over a changing population. This last point is not a defect of imputation but a restatement of the estimand problem, and it is precisely what motivates the within-cohort estimator we introduce next: rather than aggregate over a changing composition, it holds the composition fixed within each cohort and is explicit about which cohorts identify each coefficient.

Table 1: Event-study coefficients under staggered adoption with differential entry and no real pre-trends.

Relative period	Pre-treatment ($k < t^*$)						Post-treatment			
	-6	-5	-4	-3	-2	-1	0	1	2	3
<i>Panel A. Heterogeneous effects ($\tau_c = -2, -6, -10$), common baseline</i>										
Sun-Abraham	-2.30	-2.31	-1.31	-1.33	0.00	—	-6.00	-6.00	-6.00	-5.98
Stacked DiD	-2.30	-2.32	-1.31	-1.33	0.00	—	-6.00	-6.00	-6.00	-5.93
Callaway-Sant’Anna	0.00	-0.02	0.00	-0.02	0.00	—	-6.00	-6.00	-6.00	-5.98
Imputation (FEct)	0.00	-0.01	0.00	-0.01	0.00	—	-6.00	-6.00	-6.00	-5.99
<i>Panel B. Homogeneous effect ($\tau = -4$), baselines 0, 4, 8 across entry cohorts</i>										
Sun-Abraham	0.03	0.00	0.00	0.00	0.01	—	-3.99	-3.99	-4.00	-3.99
Stacked DiD	0.03	0.00	0.00	0.00	0.01	—	-3.99	-3.99	-4.00	-4.00
Callaway-Sant’Anna	-3.97	-4.00	-2.00	-2.00	0.01	—	-3.99	-3.99	-4.00	-3.99
<i>Pooled (no unit FE)</i>	-3.98	-4.01	-2.01	-2.00	0.00	—	-4.00	-4.00	-4.00	-4.00
Imputation (FEct)	0.02	-0.01	0.00	0.00	0.00	—	-4.00	-4.00	-4.00	-4.00

Notes: Panel A isolates the demeaning phantom by allowing for heterogeneous effects across entry cohorts while holding baselines common. Panel B isolates composition bias by holding effects homogeneous at $\tau = -4$ while letting baseline levels differ across cohorts. Neither setting contains a unit-level trend, and PT-W holds throughout, so every true pre-period coefficient is zero.

Table 3: Which estimators fail, and why.

	Demeaning phantom (effect het. $\times$ imbalance)	Baseline composition (baseline diff. $\times$ imbalance)	Survivor estimand (effect het. $\times$ diff. exit)
Sun–Abraham	fails	✓	phantom-distorted
Stacked DiD	fails	✓	phantom-distorted
Callaway–Sant’Anna	✓	fails	honest drift
Imputation (FEct)	✓	✓	honest drift

Notes: “Phantom” denotes the demeaning bias of Propositions 1–2. “Composition” denotes the estimand problem of Section 4. A checkmark means the estimator is robust to that channel.

Table 2: Event-study coefficients under staggered adoption with differential exit and no real dynamic effects.

Relative period	Pre			Post-treatment ($k \geq t^*$)						
	–3	–2	–1	0	1	2	3	4	5	6
<i>Panel A. Heterogeneous effects ($\tau_c = -2, -6, -10$), common baseline</i>										
Sun–Abraham	–0.01	–0.02	—	–6.01	–6.01	–6.02	–7.15	–7.15	–8.03	–8.02
Stacked DiD	–0.01	–0.02	—	–6.01	–6.01	–6.02	–7.15	–7.15	–8.03	–8.02
Callaway–Sant’Anna	–0.01	–0.02	—	–6.01	–6.01	–6.02	–8.02	–8.03	–10.03	–10.01
Imputation (FEct)	0.00	0.00	—	–6.00	–6.00	–6.01	–8.00	–8.01	–10.01	–10.00
<i>Panel B. Homogeneous effect ($\tau = -4$), baselines 0, 4, 8 across exit cohorts</i>										
Sun–Abraham	–0.01	0.00	—	–4.00	–4.00	–4.01	–4.00	–3.99	–4.02	–4.00
Callaway–Sant’Anna	–0.01	0.00	—	–4.00	–4.00	–4.01	–2.00	–1.99	–0.01	0.01
<i>Pooled (no unit FE)</i>	0.00	0.00	—	–3.99	–4.00	–4.00	–1.99	–1.98	–0.01	0.01
Imputation (FEct)	0.00	0.00	—	–4.00	–4.00	–4.00	–4.00	–3.99	–4.02	–4.00

Notes: Panel A isolates phantom dynamics by allowing for heterogeneous effects across exit cohorts while holding baselines common. Panel B isolates composition bias by holding effects homogeneous at $\tau = -4$ while letting baselines differ across cohorts. No unit’s effect is dynamic, so the true post-treatment path is flat.

6 Published Parallel-Trends Diagnostics

The preceding sections identified two problems that panel imbalance creates for conventional event studies: unit fixed effects can generate mechanical pre-treatment movement, and coefficients at different event times can describe different populations. Both affect how applied researchers evaluate parallel trends. Event-study coefficients and raw treated and control means are commonly read as evidence for pooled parallel trends (PT-P). In an imbalanced panel, successive points in these displays may describe different populations.⁹ When an event-study is estimated with unit fixed effects and treatment effects vary across cohorts, its pre-treatment path may also contain the demeaning phantom. The published path alone cannot separate that mechanical component from a genuine pre-treatment difference. We therefore examine what published diagnostics establish, without claiming to measure the phantom in any one application.

Unbalanced panels are common in difference-in-differences applications ([Fourniaies and Hall, 2022](#); [Latura and Weeks, 2023](#); [Chiu et al., 2025](#)). We examine one such case in detail, [Liao \(2023\)](#), a recent article in a leading political science journal that uses a non-staggered design with an unbalanced panel. Aligning the estimation sample with the intended estimand and constructing a more comparable control group yields a substantively different result. Appendix B reports two further replications of [Weschle \(2021\)](#) and [Clarke \(2020\)](#). In both, changes in composition across the pre-treatment periods create the same estimand problem.

6.1 Changing Composition in Parallel-Trends Diagnostics

Our primary case is [Liao \(2023\)](#), who asks whether firms that lobbied on immigration in 2017 received more favorable treatment in H-1B visa adjudication. Treatment onset is 2017 and common to all treated firms, which isolates panel imbalance from staggered adoption. Because the paper’s primary parallel-trends display is a raw-means plot, the composition

⁹Changing samples may still identify a stable population under a missingness assumption, but the conventional plot does not test one.




Figure 7: The published parallel-trends display, and the firms behind it.

Notes: Mean H-1B denial rate by calendar year for firms that did and did not lobby on immigration in 2017, reproducing Liao (2023) exactly. N_T and N_C , printed above each series, give the distinct treated and control firms behind that year’s means, whose composition need not be the same across years. No unit fixed effects. The dotted line marks treatment onset.

question arises without reference to any estimator. The panel is severely unbalanced: about a tenth of the possible firm–year cells between 1992 and 2017 are observed.

The published raw-means display. Liao plots average H-1B denial rates for firms that did and did not lobby on immigration in 2017, describes the pre-treatment trends as “quite similar,” and states that the figure increases confidence in the assumption. Figure 7 reproduces the series from the author’s replication materials and adds the firm counts behind selected points. Those counts change across the range of the data: the treated series grows from 20 firms in 1992 to 74 in 2017, and only 0.6% of the control firms observed in 2017 are present in 1992. The figure therefore does not trace a fixed set of firms through time. The firms underlying its pre-treatment years are largely not the firms behind the estimated effect.




Figure 8: The paper’s own event study, and the firms behind it.

Notes: Reproduction of Liao (2023) Appendix Table D.14, column 2: the author’s estimator (firm and year fixed effects, standard errors clustered by firm), reference period (2016) and outcome. All 25 coefficients and standard errors match the printed table. The red point is 2017. Percentages give the share of firms observed at onset in 2017 that are also present in that year, **T** treated and **C** control.

The event study. Liao also reports an event study, in Appendix Table D.14. Column 2 interacts the 2017 lobbying treatment indicator with year indicators, includes firm and year fixed effects, and clusters standard errors by firm. This is the canonical specification of Equation (1), applied to Liao’s panel. 2016 is omitted and serves as the reference period. Because 2017 is the only post-treatment year in the panel, the column contains 24 pre-treatment coefficients and a single post-treatment coefficient, which is the estimated effect of lobbying relative to 2016. We reproduce all 25 coefficients and standard errors to three decimal places.¹⁰ Figure 8 plots the corresponding event study.

The paper characterizes the pre-treatment estimates as “either small and imprecisely estimated or in the opposite direction” of the negative estimate for 2017. This characterization is accurate for each individual coefficient: of the 24 estimates, 23 are positive, and the lone

¹⁰Our reproduction runs the author’s own estimator on the author’s own data, recovering the printed observation and firm counts exactly.

negative estimate is -0.001 . Nonetheless, 10 are significant, so their sign is the only regard in which they are reassuring. Signs alone also reveal little about whether the coefficients exhibit a systematic pattern over time, which is the relevant diagnostic for assessing parallel trends.

More importantly, the composition of the sample supporting these coefficients changes substantially across the window, which Figure 8 reports year by year. The limited overlap with the onset sample demonstrates the scale of this change. Of the control firms observed at treatment onset, only 3% are also observed in 1997 and 13% in 2002. Even in 2016, the final pre-treatment year, the share is just 48%. Successive coefficients therefore do not trace the pre-treatment path of a fixed population, much less that of the firms whose 2017 outcomes produce the estimated effect.

The raw-means display and the event study both target PT-P, and both rest on firms that are largely not the ones identifying the 2017 estimate. Neither speaks to whether the treated and control firms would have followed parallel paths absent treatment. Liao’s panel is deep enough to repair the diagnostic by holding composition fixed, which is what we turn to in the following section.

6.2 Aligning the Diagnostic with a Fixed Population

We next build up a diagnostic that describes the same firms as the estimate itself. Table 4 reports the sequence. We hold the paper’s event-study estimator, outcome, reference period and treatment definition fixed throughout. We begin by changing the estimation sample and then turn to constructing a more comparable control group.

Imposing fixed composition on the published 1992–2017 window would leave only 15 of the 74 treated firms observed at onset, so the resulting diagnostic would describe a small fraction of the treated sample.¹¹ We therefore shorten the window before imposing fixed composition.

¹¹Liao (2023)’s treatment indicator is defined for the 89 treated firms of Table 4 row A, but only 74 have an observed denial rate in 2017. The remaining firms contribute pre-treatment observations without entering the 2017 contrast.

We use 2011–2017 as a practical compromise between pre-treatment history and sample retention. It provides six pre-treatment years while retaining 65 of the 74 treated firms. Appendix C reports the corresponding tradeoff across alternative starting years. Shortening the window alone leaves the estimated effect essentially unchanged, at -2.73 percentage points against -2.69 over the published span (Table 4, rows A and B).¹² Imposing fixed composition on the 2011–2017 window moves the coefficient to -2.11 percentage points (row C) and leaves six pre-treatment periods supported by a common set of firms. The shorter window therefore makes a fixed-population diagnostic feasible for most of the treated firms observed at onset, at the cost of a longer pre-treatment history.

With the population held fixed, the pre-treatment coefficients now describe the same firms over time. The resulting path, however, is not flat. The 2011 coefficient is $+2.6$ percentage points, and the estimates decline toward the 2016 reference year from there (Figure 9, left panel). A joint test that all five coefficients are zero yields $p = 0.066$ and therefore does not reject. A linear specification instead estimates a 2.10 percentage point decline in the treated-control difference between 2011 and 2016, and a test of no linear trend rejects with $p = 0.006$. The two tests ask different questions: the joint test allows the five coefficients to depart from zero independently, while the linear test concentrates power on a smooth alternative (Roth, 2022). Their results are therefore not inherently contradictory. With composition held fixed, the diagnostic tests PT-P among the firms that identify the effect. That makes it interpretable, but the evidence is not reassuring and does not establish that the full control pool provides a credible counterfactual. We turn to the choice of control group next.

¹²Liao’s headline specification (his Figure 7 and Appendix Table D.6 column 1) replaces the year interactions with a single post-treatment indicator, which measures 2017 against all pre-treatment years pooled rather than against 2016 alone. Using that specification, the four rows of Table 4 give -4.46 (the main reported estimate), -4.07 , -3.26 and -1.11 .

6.3 Choosing a Comparable Control Group

Fixed composition aligns the population but does not make the control group comparable to the treated firms. A credible comparison requires control firms that resemble the treated firms in the ways that bear on their outcomes. On the fixed-composition sample, we examine how treated and control firms differ across a set of firm-level covariates, which we pull from Orbis in line with [Liao \(2023\)](#). The standardized differences reach 1.5 on size class and 1.6 on public listing, and the entry-cohort distributions overlap only partially (Table 7). Matching serves three purposes here: it absorbs unobserved differences associated with entry cohort, it builds a control group comparable on the covariates the paper itself uses, and it improves the pre-treatment comparison.

We therefore replace the full control pool with firms matched to the treated firms within entry cohort. Matching within cohort eliminates differences attributable to cohort, including unobserved cohort-level characteristics associated with the length of a firm’s filing history. A firm that first appears in 1992 has been filing H-1B petitions for two decades before our estimation window opens and is unlikely to be the same kind of firm as one that first appears in 2011. Each treated firm is matched to two controls without replacement on a pooled propensity score combining the 2011–2015 denial rates and petition volume with the Orbis industry, size class and listing indicators.¹³ [Liao \(2023\)](#) matches on those covariates in a two-period specification ([Imai et al., 2023](#)) that compares 2016 with 2017 alone and uses no pre-treatment outcomes, so it cannot speak to whether parallel trends holds.

Of the 65 fixed-composition treated firms, 64 merge to the required Orbis covariates, and all 64 are matched to 128 control firms. Cohort composition is identical by construction, the standardized difference on size class falls from 1.5 to 0.04, and the total absolute difference in industry shares falls from 1.0 to 0.22. Outcome balance improves as well, though the 2011–2015 denial rates enter the score, so their standardized difference falling from 0.515 to

¹³Entry cohort also enters the pooled score and is the only exact-matching constraint, ensuring that no match crosses cohorts.




Figure 9: The paper’s own specification on a fixed-composition sample, before and after matching within entry cohort.

Notes: Event-study interactions in the paper’s own specification, with firm and year fixed effects, standard errors clustered by firm and 2016 omitted, unchanged between panels. The dotted line marks treatment onset. The left panel is row C of Table 4 and the right is row D, both on firms observed continuously from 2011 through 2017. The annotation gives the joint test of the five pre-treatment coefficients.

0.019 reflects that fit rather than testing it. The 2016 reference year is the informative case, because we withhold it from the score, and its standardized difference falls from 0.378 to 0.134.

Matching also attenuates the apparent pre-treatment trend. With the full control pool, the fitted trend implies a 2.10 percentage point decline over 2011–2016. With the matched controls, the implied decline is 1.05 percentage points. The joint-test p -value rises from 0.066 to 0.719, and the linear-trend p -value rises from 0.006 to 0.258.¹⁴ The matched coefficients nevertheless remain roughly one percentage point above the reference (Figure 9). We therefore interpret the matching as improving observed comparability, not as establishing that either PT-P or PT-W holds.

With comparability improved, the estimated effect in the matched fixed-composition sample is -0.43 percentage points, with a standard error of 1.16 (Table 4, row D). The 95%

¹⁴For the linear pre-trend test, a wild cluster bootstrap gives $p = 0.270$, compared with an analytic $p = 0.258$ (Cameron et al., 2008). Both are conditional on the selected match.

Table 4: Sample composition and the estimated effect at each step.

Sample	Firms		Observed	2017
	Treated	Control	throughout	vs. 2016
A 1992–2017, as published	89	384,373	17% / <0.1%	−2.69 (0.63)
B 2011–2017 window	86	126,786	76% / 5.0%	−2.73 (0.63)
C 2011–2017 (fixed composition)	65	6,303	100% / 100%	−2.11 (0.69)
D + matched within cohort	64	128	100% / 100%	−0.43 (1.16)

Notes: Estimated effect in percentage points, with standard errors clustered by firm in parentheses. Every row uses the specification of Liao (2023, Appendix Table D.14, column 2): firm and year fixed effects, the paper’s outcome and treatment definitions, and 2016 as the reference period. Row A reproduces that column’s 2017 coefficient. “Observed throughout” gives the share of a row’s firms present in every year of its window, treated share first. Row D matches 2 control firms to each treated firm within entry cohort, on the propensity score of Section 6.3. The Orbis merge costs 1 treated firm.

confidence interval is wide enough to include zero. Our estimate also concerns a narrower population: the 64 treated firms observed in every year from 2011 through 2017 that merge to the Orbis covariates. The estimate is noisier, but the quantity it targets is better identified: the pre-treatment diagnostic and the estimated effect rest on the same firms.

In an unbalanced panel, the units used to estimate pre-treatment event-study coefficients can vary across periods and differ from those identifying the post-treatment effect. Therefore, they do not provide direct evidence on how the treated and control units identifying the effect were trending prior to treatment. In Liao’s panel, shortening the window leaves enough data to hold composition fixed while retaining six pre-treatment periods, and matching within entry cohort supplies a more comparable control group. Together, these changes produce a substantially different result. In other settings, where fixed composition would instead eliminate most of the treated units, subsetting cannot provide the same relief. Section 7 proposes a within-cohort estimator for such settings.

7 A Within-Cohort Estimator

The preceding sections establish the phantom pre-trend result analytically and reproduce it in simulation. Section 6.2 repaired Liao (2023)’s diagnostic by subsetting, because a seven-year window with fixed composition still retained 65 of the 74 treated firms observed at

onset.

Not every panel is that forgiving. Unbalanced panels in which fixed composition would remove most of the treated units are not unusual ([Clarke, 2020](#); [Weschle, 2021](#)), and our own application is one: requiring defectors to be observed throughout the displayed pre-period would leave 6 of 23.¹⁵ Where subsetting is that costly, the requirements of Section 4 have to be built into the estimator. This section develops an estimator that does so, holding entry cohorts fixed rather than discarding the units whose entry dates differ, and applies it to the 2015 revolt against Speaker Boehner.

7.1 The Empirical Setting

The U.S. House elects its Speaker by floor vote at the opening of each Congress, almost always along party lines. At the opening of the 114th Congress in January 2015, 29 Republicans declined to support the sitting Speaker, John Boehner, which was a striking defection in a vote that is rarely contested. Party discipline is sustained in part by leaders’ ability to reward loyalty and punish defection. One plausible lever of punishment runs through fundraising. Leadership is closely tied to a network of corporate donors, and defectors may find that access to those donors is subsequently curtailed ([Canes-Wrone and Miller, 2022](#); [Heberlig et al., 2006](#); [Kates and Thieme, 2023](#); [Issue One, 2017](#); [Martin, 2021](#)). Whether such punishment actually materializes has not, to our knowledge, been established empirically. We therefore ask: what is the effect of defecting in the Speaker election on a defector’s subsequent fundraising from the corporate donors connected to party leadership?

As in [Liao \(2023\)](#), the design is non-staggered and unbalanced. All defectors receive treatment at the common onset of the 114th Congress ($t^* = 114$). The defectors entered the House across five Congresses, from the 110th through the 114th, and therefore form five entry cohorts. We retain 23 of the 29 defectors, dropping the six who did not run for reelection following the revolt and whose post-treatment fundraising is not observed. The

¹⁵[Clarke \(2020\)](#) studies faction membership among U.S. House candidates, [Weschle \(2021\)](#) outside earnings among UK MPs. Appendix B discusses both.

control pool consists of Republican House members who did not defect.

The outcome Y_{it} is the log dollar amount that legislator i raises in Congress t from corporate donors connected to party leadership, constructed from the Database on Ideology, Money in Politics, and Elections (Bonica, 2024). We classify a corporate donor as leadership-connected when it also contributes to the legislator’s own-party leadership in the same Congress, and use $\log(\text{dollars} + 1)$ to accommodate zeros.

Entry timing determines the length of pre-treatment history each legislator contributes. A member first elected to the 110th Congress is observed as an incumbent for four pre-treatment Congresses, from the 110th through the 113th, whereas a member entering in the 114th has no incumbent pre-period. For members entering in the 113th and 114th Congresses, we use fundraising from their challenger campaign as an additional pre-treatment observation. This supplies the 113th-Congress reference value for the 114th cohort and an earlier, 112th-Congress value for the 113th cohort. Earlier cohorts nevertheless contribute longer pre-treatment histories: $T_{\text{pre}}(c)$ declines across later cohorts while the post-period is shared, which is the differential-entry structure studied in Section 2.2.

Two features of the setting bear on what its pooled diagnostic can show. First, the panel is imbalanced and the estimated effects of defection differ across entry cohorts (Table 5, below), which is the condition under which Propositions 1–2 predict a phantom pre-trend. The pooled coefficients therefore carry a phantom component added to whatever genuine pre-treatment difference is present (Remark 1). Nothing in these data separates the two. Second, unlike the stylized DGP of Section 3.1, parallel trends is not guaranteed to hold. The question is therefore whether it can be made credible, by the within-cohort matching that improved the comparison for Liao in Section 6.3.

We proceed in three steps. We begin with the pooled event study, then hold entry cohort fixed, and finally match defectors to comparable controls within cohort. The last step is the estimator this paper proposes, which Section 7.3 states in general form before applying it here.

7.2 The Pooled Event Study

We begin with the pooled event study of Equation (1), estimated on all 23 defectors against the full pool of non-defecting Republicans, with legislator and Congress fixed effects and the 113th Congress omitted as the reference period. Figure 10 reports it together with the legislators behind each coefficient.




Figure 10: An apparent pre-trend against the full control pool, and the legislators behind it.

Notes: Pooled event-study coefficients from Equation (1), with the 113th Congress omitted and the dotted line marking onset in the 114th. Counts give the distinct legislators behind each coefficient, **T** treated and **C** control. Bars are 95% confidence intervals with standard errors clustered by legislator.

The pre-treatment coefficients are large and individually distinguishable from zero, at 0.65, 0.55 and 0.57 log points in the 110th through the 112th Congresses. Because the 113th Congress is the reference, the coefficients indicate that the earlier treated–control gaps were substantially larger than the gap immediately before the revolt. Read conventionally, this pattern would be taken as evidence against parallel trends and could lead a careful researcher to abandon the design. Yet it raises two distinct concerns. The full control pool may not provide a credible comparison for the defectors, and the pre-treatment coefficients

may describe different populations.

First, the full pool of non-defecting Republicans provides a weak counterfactual for the defectors. The two groups differ substantially in their pre-treatment fundraising and their relative gap narrows across the observed pre-period, so pooled parallel trends (PT-P) fails against this control group. Whether within-cohort parallel trends (PT-W) can be made to hold against better-chosen controls is the question Section 7.3 takes up.

Second, the pre-treatment coefficients describe different sets of defectors, which is the estimand problem of Section 4 and Liao’s diagnostic (Figure 7). The counts in Figure 10 give the scale. The 110th-Congress coefficient is identified from the six defectors who had entered the House by that Congress, the 111th adds the single defector entering in that cohort, and the treated sample reaches all 23 only in the 114th. The earliest coefficients therefore describe a quarter of the defectors who ultimately identify the treatment effect. They neither test parallel trends for the full treated population nor remain directly comparable as later cohorts enter.

To separate these concerns, we next estimate the event study separately by entry cohort, continuing to use all non-defecting Republicans in each cohort as controls. We omit legislator fixed effects and end the panel in the 114th Congress, the last period in which every included member is observed. The cohort restriction prevents the coefficients from pooling legislators with different entry histories, while the common endpoint holds composition fixed within cohort. Without legislator fixed effects, the coefficients contain no demeaning and therefore no phantom component. Even so, the top panel of Figure 11 provides unfavorable pre-treatment evidence. The joint pre-treatment test rejects in three of the four testable cohorts, with p -values of 0.042, 0.003 and 0.016. The fourth approaches rejection at $p = 0.055$. The pooled result is therefore not only a consequence of combining incomparable cohorts. Even within cohort, the full control pool does not reproduce the defectors’ pre-treatment fundraising. This remaining comparison problem motivates matching.




Figure 11: Matching within cohort largely removes the within-cohort pre-trend.

Notes: Cohort-specific event studies estimated without legislator fixed effects. The top panel uses all non-defecting Republicans in each entry cohort, the bottom uses 1:2 within-cohort propensity-score matches. Each panel reports at bottom left the legislators behind its cohort, **T** treated and **C** control, and a joint pre-treatment test. Counts are fixed across the displayed periods, and the change in **T** between the panels is due to within-cohort trimming. The 113th Congress is omitted. The 114th cohort has no testable coefficient before onset. Bars are 95% confidence intervals with standard errors clustered by legislator.

7.3 A Within-Cohort Matching Estimator

We address this remaining comparison problem by matching each treated unit to controls from the same entry cohort on pre-treatment outcome trajectories and covariates. We then estimate a cohort-specific average treatment effect on the treated without unit fixed effects and aggregate the estimates across cohorts. The matching step resembles PanelMatch (Imai et al., 2023), but the two procedures define the pool of eligible control observations differently. PanelMatch matches on treatment history. In our non-staggered setting with common

treatment history, we instead match exactly on entry cohort. Because comparisons do not cross cohorts, baseline differences across cohorts no longer motivate unit fixed effects. We can therefore omit them and avoid the demeaning that produces phantom pre-trends.¹⁶ The procedure has three steps.

1. **Define cohorts.** Group units by common observation history. When all units remain observed through a common final period and have no interior gaps, entry alone determines that history and T_c is constant within cohort. This is the setting defined in Section 2.1. If units exit at different times, define cohorts jointly by entry and exit. Units with interior gaps should be dropped, as their histories are often too idiosyncratic to pool.
2. **Match within cohort.** Match each treated unit to controls drawn only from its own cohort using pre-treatment outcome trajectories and covariates. Available methods include nearest-neighbor, Mahalanobis, propensity-score, and coarsened exact matching. Weighting and outcome modeling provide alternatives to matching, including augmented inverse-probability weighting for doubly robust estimation. Remove treated units whose pre-treatment trajectories fall outside the support of the controls in their own cohort, since the observed data contain no comparable counterfactual for them.
3. **Estimate, aggregate and conduct inference.** Estimate a cohort-specific effect on the matched sample without unit fixed effects. Aggregate the estimates with weights proportional to the number of retained treated units $n_{1,c}$. Base inference on the full procedure rather than on the final regression alone.

Under stronger assumptions that may be credible in more data-rich settings, one can model individual-unit heterogeneity by requiring parallel trends to hold for each treated unit and its matched control, rather than only in expectation within a matched cohort.

¹⁶Each cohort is a balanced sub-panel, so the canonical event study could be estimated within cohort. Once matching has been done, including unit fixed effects is inefficient.

Implementation. The 23 defectors entered the House across five Congresses. None is absent from a Congress between entering and the revolt so entry alone therefore defines the five cohorts. Truncating the panel at the 114th Congress imposes the common endpoint assumed in Section 2.1. Challenger fundraising supplies the 113th-Congress value for the three defectors entering in the 114th. Trimming for within-cohort support removes two defectors from the 112th cohort and one from the 113th, leaving 20 treated legislators. Appendix Figure 14 shows the candidate-level observations underlying this restriction.

To construct the control groups, we estimate a single logistic propensity-score model across cohorts using the available pre-treatment log outcomes, immediate pre-period ideology (a dynamic CF score), and district vote share.¹⁷ Pooling the model allows the covariate coefficients to draw information from all five cohorts, while exact matching on cohort prevents comparisons from crossing them. Each retained defector is matched without replacement to two non-defecting Republicans using nearest-neighbor matching on the estimated score. Within each matched cohort, we estimate a 2×2 difference-in-differences comparing the 113th and 114th Congresses. For the 114th entry cohort, challenger fundraising supplies the 113th-Congress value. We aggregate the five estimates using weights proportional to the number of retained treated units $n_{1,c}$ and bootstrap the full matching-and-estimation procedure.

Differential exit. The main analysis handles differential exit by ending the panel in the 114th Congress. This preserves fixed composition but limits the estimate to the treatment-onset period. Four defectors leave the House after the revolt, so extending the cohort event studies would compare a shrinking treated population and reintroduce the composition problem.

Later post-treatment outcomes can still be informative if the change in target population is made explicit. One option is to average the Congresses after the 114th into a

¹⁷Each pre-treatment Congress enters through an indicator for whether the legislator is observed and its interaction with the corresponding log outcome, with missing outcomes set to zero. The challenger periods enter analogously.

single unit-level observation. This approach gives every legislator observed in at least one later Congress a single post-treatment observation, preventing longer-serving members from receiving greater weight simply because they contribute more rows. Appendix Figure 16 extends the cohort event studies one period further using this construction. However, pooling does not eliminate every consequence of exit. Four defectors have no post-114 outcome at all, so the additional point describes a smaller population than the treatment-onset estimate. Because the extension omits unit fixed effects, it contains no demeaning phantom, but the change in population across post-treatment estimates should be reported alongside them.

Pre-treatment diagnostics. The bottom panel of Figure 11 repeats the cohort-specific event studies after support trimming and within-cohort matching. Relative to the full-sample estimates against the full control pool in the top panel, the pre-treatment coefficients move closer to zero. The joint tests do not reject in any of the four testable cohorts, with the smallest p -value equal to 0.09 in the 113th cohort. The 114th cohort has only the 113th-Congress reference observation before treatment, so its pre-treatment path cannot be assessed. The contrast between panels therefore shows that support trimming and matched control selection substantially improve the observed within-cohort comparison.

These diagnostics do not provide an independent validation of PT-W. The pre-treatment outcomes underlying the displayed coefficients also enter the propensity score, so the improved fit is partly a product of the procedure used to select controls. The results show that the selected controls reproduce the defectors' observed pre-treatment histories more closely, not that their counterfactual trends would have remained parallel after treatment. Appendix Figure 15 reports the corresponding raw means, and Appendix Figure 17 reports balance on ideology and district vote share.

Inference. For the aggregate estimate, we bootstrap the full matching-and-estimation procedure rather than calculate a standard error conditional on the realized match. Each replicate resamples the control pool, repeats the matching and cohort-specific estimation, and

aggregates the resulting estimates. The treated legislators remain fixed across replicates. Inference therefore incorporates uncertainty from selecting the controls while remaining conditional on the 20 defectors retained after trimming, which are the treated population for the estimand.

Results. Table 5 reports the cohort-specific and aggregate estimates. Among the 20 defectors retained after trimming, the aggregate estimate is a 0.46 log-point reduction in leadership-connected corporate fundraising (bootstrap SE 0.057, 95% CI $[-0.62, -0.39]$). This implies that defectors raised approximately 37% less from these donors following the revolt.

The cohort estimates range from -0.15 to -0.66 log points. Because the cohorts are small, with only one retained defector in the 111th cohort, we treat this variation as descriptive rather than as evidence of statistically significant differences across cohorts. The aggregate estimate applies to the 20 defectors in the trimmed, matched sample, not to all 29 Republicans who defected.

Table 5: Cohort-specific and aggregate effects of defection on corporate fundraising.

Entry cohort (Congress)	Defectors $n_{1,c}$		
	All	After trimming	ATT (log scale)
110	6	6	-0.47
111	1	1	-0.15
112	8	6	-0.66
113	5	4	-0.45
114	3	3	-0.17
Total	23	20	-0.46

Notes: Rows are entry cohorts, defined by the Congress in which a defector first entered the House, so earlier cohorts are observed for more pre-treatment periods $T_{\text{pre}}(c)$. The first count column gives all defectors in the cohort and the second those remaining after the within-cohort trimming of Section 7.3. Cohort ATTs come from within-cohort 1:2 propensity-score matching, and the aggregate weights cohorts by their trimmed $n_{1,c}$.

What to report. We recommend three reporting practices for within-cohort designs. First, define the target population explicitly. If a cohort has too little pre-treatment history to support a diagnostic, report that limitation rather than borrowing reassurance from other cohorts. Second, report treated and control counts for every displayed period together with their overlap with the treatment-period population. Counts alone can conceal substantial turnover when group size remains stable. Third, when summarizing cohort-specific diagnostics, use a clearly defined common-composition sample and fixed weights. Report the disaggregated paths alongside the aggregate.

Staggered adoption. The same logic extends to staggered adoption, but cohort definitions must incorporate both observation history and adoption timing. Under differential entry, this produces entry-by-adoption-time cells. If exit also varies, exit timing must enter the definition as well. Section 5 shows why estimators robust to heterogeneous treatment timing do not necessarily solve the imbalance problem. They partition the data by adoption timing rather than by entry or exit, leaving imbalance within each adoption cohort. For estimators that demean within those cells, the T_{post}/T_c mechanism remains unchanged. The number of cells grows quickly as these dimensions accumulate, so cell size is the practical constraint on the extension.

In this application, the pooled and within-cohort diagnostics point in opposite directions. The pooled display combines a comparison problem with a population problem and separates neither. Analyzing each cohort separately fixes the population being compared, while matching improves the choice of controls within it. The resulting estimate comes from a design that a conventional reading of Figure 10 would have discarded.

8 Discussion

In recent years, numerous papers have demonstrated issues with two-way fixed effect designs under effect-heterogeneity. However, this work has primarily focused on settings with

staggered treatment-adoption, treating the common-onset setting (where ever-treated units enter treatment in the same period) as safe, and typically assumes a balanced panel in their theoretical setup, despite settings where units are observed for different sets of periods being quite common.

This paper shows that even in the common treatment-onset setting, panel imbalance creates two distinct problems for the canonical event-study estimator. First, we prove that the estimator is biased under effect-heterogeneity. When units enter (exit) the sample in different periods, the estimator may fabricate non-existent, or conceal genuine, pre-trends (dynamic effects). Second, even setting the bias aside, the estimator targets the wrong estimand for assessing identification and dynamic effects. Coefficients may be estimated on a different set of units than those used to estimate the treatment effect. Therefore, pre-treatment coefficients for periods before all units have entered do not provide evidence of how the units identifying the ATT were trending prior to treatment. And analogously, post-treatment coefficients after some units have exited are not directly comparable to the treatment effect. Researchers may conflate the changing composition of units with evidence against or for pre-trends or dynamic effects.

To resolve both problems, we define effects within cohorts of units observed in the same periods. This holds composition fixed within cohort, permits comparisons against cohort-specific controls, and avoids the demeaning that produces the phantom.

The practical implications can vary sharply across applications. A pooled pre-trend that would ordinarily end a project can overstate the problem, and a flat pooled path can hide one. In the Boehner application, the pooled event study reports large, positive pre-treatment coefficients. Read conventionally, that display would lead a careful researcher to abandon the design. After support trimming and matching within entry cohort, the joint pre-treatment tests no longer reject in any of the four testable cohorts. The aggregate estimate is a substantial reduction in leadership-connected corporate fundraising. In [Liao \(2023\)](#)'s panel, the published display invites the opposite conclusion. Its apparently similar

pre-treatment paths describe a changing population, while holding composition fixed reveals a declining path that is difficult to reconcile with the reassuring interpretation. Matching controls within entry cohort further improves the comparison and attenuates the estimated effect to a small, imprecise difference that includes zero. The pooled displays therefore obscure different problems in the two applications, rejecting a design that becomes more credible within cohort in one and providing reassurance for the wrong population in the other.

The applications also illustrate two responses to changing composition. When the data permit, researchers can hold the diagnostic population fixed directly. In Liao’s panel, shortening the window makes this possible while preserving most of the treated firms observed at onset. In the Boehner application, imposing fixed composition over the available pre-treatment span would instead discard most of the defectors. We therefore estimate within entry cohort, preserving most of the treated sample after support trimming. In both applications, the full control pool fails to reproduce the treated units’ pre-treatment paths once the diagnostic is aligned with the relevant population. Matching within cohort improves the comparison, but the resulting diagnostics demonstrate observed comparability rather than establish PT-W.

Three extensions merit further work. First, our formal treatment emphasizes differential entry. Differential exit changes the post-treatment share of the periods used in demeaning and can produce apparent dynamics even when no unit’s effect changes. The exit case warrants the same formal development. Second, under staggered adoption, cohorts must incorporate adoption timing as well as entry and exit. Section 5 shows that the mechanism survives within adoption cohorts whose members have different observation histories. The practical limitation is that these cells thin quickly as the number of dimensions grows. Third, matching is only one implementation of the within-cohort framework. Weighting, outcome modeling, and doubly robust estimators can also be applied within cohort. Developing inference and sensitivity analyses for these alternatives would strengthen the approach where

overlap is limited or the adjustment model is uncertain.

The central lesson is that parallel trends is not a property of an event-study plot in isolation. It is a claim about the counterfactual path of a defined population against a defined comparison group. In an unbalanced panel, a useful diagnostic must make both explicit.

References

- Acemoglu, Daron, Suresh Naidu, Pascual Restrepo, and James A. Robinson. 2019. “Democracy Does Cause Growth.” *Journal of Political Economy*, 127(1): 47–100.
- Angrist, Joshua D., and Jörn-Steffen Pischke. 2009. *Mostly Harmless Econometrics: An Empiricist’s Companion*. Princeton: Princeton University Press. (§5.2, regression DiD with leads and lags.)
- Autor, David H. 2003. “Outsourcing at Will: The Contribution of Unjust Dismissal Doctrine to the Growth of Employment Outsourcing.” *Journal of Labor Economics*, 21(1): 1–42.
- Baker, Andrew C., David F. Larcker, and Charles C. Y. Wang. 2022. “How Much Should We Trust Staggered Difference-in-Differences Estimates?” *Journal of Financial Economics*, 144(2): 370–395.
- Bonica, Adam. 2024. Database on Ideology, Money in Politics, and Elections: Public version 4.0 [Computer file]. Stanford, CA: Stanford University Libraries. <https://data.stanford.edu/dime>.
- Borusyak, Kirill, Xavier Jaravel, and Jann Spiess. 2024. “Revisiting Event-Study Designs: Robust and Efficient Estimation.” *Review of Economic Studies*, 91(6): 3253–3285.
- Callaway, Brantly, and Pedro H. C. Sant’Anna. 2021. “Difference-in-Differences with Multiple Time Periods.” *Journal of Econometrics*, 225(2): 200–230.
- Cameron, A. Colin, Jonah B. Gelbach, and Douglas L. Miller. 2008. “Bootstrap-Based Improvements for Inference with Clustered Errors.” *Review of Economics and Statistics*, 90(3): 414–427.
- Canes-Wrone, Brandice, and Kenneth M. Miller. 2022. “Out-of-District Donors and Representation in the US House.” *Legislative Studies Quarterly*, 47(2): 361–395.

- Clarke, Andrew J. 2020. "Party Sub-Brands and American Party Factions." *American Journal of Political Science*, 64(3): 452–470.
- Chiu, Albert, Xingchen Lan, Ziyi Liu, and Yiqing Xu. 2025. "Causal Panel Analysis under Parallel Trends: Lessons from a Large Reanalysis Study." *American Political Science Review*. Forthcoming.
- de Chaisemartin, Clément, and Xavier D'Haultfœuille. 2020. "Two-Way Fixed Effects Estimators with Heterogeneous Treatment Effects." *American Economic Review*, 110(9): 2964–2996.
- Fourinaies, Alexander, and Andrew B. Hall. 2022. "How Do Electoral Incentives Affect Legislator Behavior? Evidence from U.S. State Legislatures." *American Political Science Review*, 116(2): 662–676.
- Goodman-Bacon, Andrew. 2021. "Difference-in-Differences with Variation in Treatment Timing." *Journal of Econometrics*, 225(2): 254–277.
- Heberlig, Eric, Marc Hetherington, and Bruce Larson. 2006. "The Price of Leadership: Campaign Money and the Polarization of Congressional Parties." *The Journal of Politics*, 68(4): 992–1005.
- Imai, Kosuke, In Song Kim, and Erik H. Wang. 2023. "Matching Methods for Causal Inference with Time-Series Cross-Sectional Data." *American Journal of Political Science*, 67(3): 587–605.
- Issue One. 2017. "The Price of Power: How Political Parties Squeeze Influential Lawmakers to Boost their Campaign Coeffers." Technical Report, Issue One.
- Jacobson, Louis S., Robert J. LaLonde, and Daniel G. Sullivan. 1993. "Earnings Losses of Displaced Workers." *American Economic Review*, 83(4): 685–709.

- Kates, Sean, and Sebastian Thieme. 2023. “Fundraising Events and Non-Ideological Donation Motivations.” IAST Working Paper No. 159, Institute for Advanced Study in Toulouse. <https://publications.ut-capitole.fr/id/eprint/48740/>.
- Latura, Audrey, and Ana Catalano Weeks. 2023. “Corporate Board Quotas and Gender Equality Policies in the Workplace.” *American Journal of Political Science*, 67(3): 606–622.
- Liao, Steven. 2023. “The Effect of Firm Lobbying on High-Skilled Visa Adjudication.” *The Journal of Politics*, 85(4): 1416–1429.
- Liu, Licheng, Ye Wang, and Yiqing Xu. 2024. “A Practical Guide to Counterfactual Estimators for Causal Inference with Time-Series Cross-Sectional Data.” *American Journal of Political Science*, 68(1): 160–176.
- Liu, Ziyi. 2025. “Cohort-Anchored Robust Inference for Event-Study with Staggered Adoption.” Working paper, arXiv:2509.01829.
- Martin, Jonathan. 2021. “Liz Cheney’s Consultants Are Given an Ultimatum: Drop Her, or Be Dropped.” *The New York Times*, October 21. <https://www.nytimes.com/2021/10/21/us/politics/liz-cheney-kevin-mccarthy.html>.
- Paglayan, Agustina S. 2021. “The Non-Democratic Roots of Mass Education: Evidence from 200 Years.” *American Political Science Review*, 115(1): 179–198.
- Roth, Jonathan. 2022. “Pretest with Caution: Event-Study Estimates after Testing for Parallel Trends.” *American Economic Review: Insights*, 4(3): 305–322.
- Sun, Liyang, and Sarah Abraham. 2021. “Estimating Dynamic Treatment Effects in Event Studies with Heterogeneous Treatment Effects.” *Journal of Econometrics*, 225(2): 175–199.

Weschle, Simon. 2021. "Parliamentary Positions and Politicians' Private Sector Earnings: Evidence from the UK House of Commons." *The Journal of Politics*, 83(2): 706–721.

A Proof of Proposition 2: the Exact Phantom Coefficients

We derive the closed form for the pre-treatment coefficients $\hat{\beta}_k$, $k < t^*$, under the data-generating process of Equation (2). Throughout, $r \equiv t^* - 1$ is the omitted reference period, $N \equiv \sum_c n_c$ the number of treated units, $\bar{\tau} \equiv \frac{1}{N} \sum_c n_c \tau_c$ their count-weighted average effect, and $\delta_c \equiv \tau_c - \bar{\tau}$.

Step 1: reduction to a treated-only regression. By the decomposition in Corollary 1, $\hat{\beta}_k = \beta_k + b_k$, where $\beta_k = \hat{\beta}_k[Y(0)]$ and $b_k = \hat{\beta}_k[W]$ with treatment component $W_{it} \equiv D_i \mathbf{1}(t \geq t^*) \tau_{c(i)}$. Under (2) the untreated potential outcome equals μ_1 for every treated unit-period and μ_0 for every control unit-period — constant within each unit — so it is absorbed entirely by the unit fixed effects and $\beta_k = 0$. Hence $\hat{\beta}_k = b_k$: the phantom is generated solely by W , which is 0 for all controls and all treated pre-periods and τ_c for treated post-periods.

The controls contribute only a normalization. Reparametrize the treated period effects as $\phi_t \equiv \lambda_t + \beta_t$ (with $\beta_r = 0$, so $\phi_r = \lambda_r$); the fitted values are $\alpha_i + \lambda_t$ for a control and $\alpha_i + \phi_t$ for a treated unit, so the two groups share a parameter only through λ_t . For $t \neq r$ the treated period effect ϕ_t is free, so treated observations impose no restriction on λ_t ; only controls do. Controls have $W \equiv 0$, so on the (connected) control subpanel the additive model fits exactly with $\hat{\lambda}_t$ constant across t ; normalize this constant to 0. Then $\hat{\beta}_k = \hat{\phi}_k - \hat{\phi}_r$, where the treated period effects $\{\phi_t\}$ and treated cohort effects $\{a_c\}$ (common within a cohort, as treated units in a cohort are identical) solve the *treated-only* two-way least-squares problem

$$\min_{\{a_c\}, \{\phi_t\}} \sum_c n_c \left[\sum_{t=e_c}^{t^*-1} (0 - a_c - \phi_t)^2 + \sum_{t=t^*}^{\tau} (\tau_c - a_c - \phi_t)^2 \right]. \quad (10)$$

Step 2: solving (10). All post-periods carry identical data ($W = \tau_c$, every cohort present), so by uniqueness $\hat{\phi}_t$ is common across $t \geq t^*$; write it p and collapse the post block to a single period of weight T_{post} . Let $S_k \equiv \{c : e_c \leq k\}$ be the cohorts observed at pre-period

k , $N_k \equiv \sum_{c \in S_k} n_c$, and write $\bar{a} \equiv \frac{1}{N} \sum_c n_c a_c$ and $\bar{a}_{S_k} \equiv \frac{1}{N_k} \sum_{c \in S_k} n_c a_c$. The first-order conditions are

$$\text{(pre-period } t) \quad \phi_t = -\bar{a}_{S_t}, \quad (11)$$

$$\text{(post)} \quad p = \bar{\tau} - \bar{a}, \quad (12)$$

$$\text{(cohort } c) \quad T_c a_c = T_{\text{post}} \tau_c - \sum_{t=e_c}^{t^*-1} \phi_t - T_{\text{post}} p. \quad (13)$$

Every cohort is present at the reference period ($e_c < t^*$, so $e_c \leq r$), hence $\hat{\phi}_r = -\bar{a}$. Combining with (11),

$$\hat{\beta}_k = \hat{\phi}_k - \hat{\phi}_r = \bar{a} - \bar{a}_{S_k}, \quad k < t^*. \quad (14)$$

The phantom coefficient at pre-period k is thus the gap between the average treated cohort effect over the full sample and its average over only the cohorts observed at k . Substituting (11)–(12) into (13) gives the linear system that pins down $\{a_c\}$:

$$T_c a_c = T_{\text{post}} \delta_c + T_{\text{post}} \bar{a} + \sum_{t=e_c}^{t^*-1} \bar{a}_{S_t}. \quad (15)$$

Step 3: proof of Proposition 2. Order cohorts by entry, $e_1 < \dots < e_C$, and suppose $\hat{\beta}_k = 0$ for all $k < t^*$; by (14) this means $\bar{a}_{S_k} = \bar{a}$ for every pre-period k . At $k = e_1$, $S_{e_1} = \{1\}$, so $a_1 = \bar{a}$. For $j = 2, \dots, C-1$, the set $S_{e_j} = S_{e_{j-1}} \cup \{j\}$ averages $\bar{a}$ having added cohort j to a set already averaging $\bar{a}$, which forces $a_j = \bar{a}$; then $a_C = \bar{a}$ follows from $\sum_c n_c a_c = N\bar{a}$. With all $a_c = \bar{a}$, the sum in (15) equals $T_{\text{pre}}(c) \bar{a}$, so (15) reduces to $T_c \bar{a} = T_{\text{post}} \delta_c + T_{\text{post}} \bar{a} + T_{\text{pre}}(c) \bar{a} = T_{\text{post}} \delta_c + T_c \bar{a}$. Hence $T_{\text{post}} \delta_c = 0$, and since $T_{\text{post}} \geq 1$, $\delta_c = 0$ for all c : effects are homogeneous. Contrapositively, if effects are heterogeneous then $\hat{\beta}_k \neq 0$ for some $k < t^*$. $\square$

Closed form with two cohorts. When $C = 2$, solving (15) gives $a_2 - a_1 = \frac{T_{\text{post}}}{T_2}(\tau_2 - \tau_1)$, so by (14) the coefficient at any pre-period observing only the earlier-entering cohort is

$$\hat{\beta}_k = \frac{n_2}{N} \frac{T_{\text{post}}}{T_2} (\tau_2 - \tau_1), \quad (16)$$

while $\hat{\beta}_k = 0$ at pre-periods observing both. The phantom is nonzero precisely when $\tau_1 \neq \tau_2$, and grows with the effect gap, the later cohort's share n_2/N , and the fraction T_{post}/T_2 of that cohort's panel lying after treatment.

Example (the design of Figure 2). Three equal-sized cohorts observed for $T_{\text{pre}} = 3, 2, 1$ pre-periods (so $T_c = 6, 5, 4$, $T_{\text{post}} = 3$) with $\tau = (-10, -6, -2)$, hence $\bar{\tau} = -6$ and $\delta = (-4, 0, 4)$. Solving (15) gives, centering $\bar{a} = 0$, the cohort effects $a = (-2.7, -0.3, 3.0)$, and (14) yields

$$\hat{\beta}_{-3} = \bar{a} - a_1 = 2.7, \quad \hat{\beta}_{-2} = \bar{a} - \frac{a_1 + a_2}{2} = 1.5, \quad \hat{\beta}_{-1} = 0,$$

matching the simulated coefficients of Figure 2 exactly.

Extensions. Two common features of applied panels leave the closed form (14) unchanged. First, a *common time trend*: adding a period effect γ_t shared by treated and control units to (2) is absorbed exactly by the period fixed effects, so it drops out of the treated-only problem (10) and the phantom is unaffected. Second, *cross-cohort baseline differences*: allowing $Y(0)$ to take cohort-specific levels — so that pooled parallel trends (PT-P) can fail even while within-cohort parallel trends (PT-W) holds — also leaves (14) intact, because each unit's level is absorbed by its own fixed effect. Controls still pin $\hat{\lambda}_t$ to a constant, and the treated demeaned outcomes, and hence the phantom, depend only on the effect heterogeneity $\{\tau_c\}$, never on the baselines.

B Further Published Designs with Panel Imbalance

This appendix reports two further published designs in which the evidence available to assess parallel trends rests on a population different from the one the estimate concerns. Both were reproduced from the authors’ own replication packages before any of the counts below were computed. As with Liao, the claim is about what a pre-treatment comparison can establish.

Weschle (2021): composition change that differs sharply across panels. [Weschle \(2021\)](#) asks whether taking up an influential position in the UK House of Commons raises an MP’s private-sector earnings. Table 1 reproduces exactly. The parallel-trends evidence is Figure 3, whose left column plots average logged earnings in the two years before taking up a position, the year of assuming it, and the first two years in it, for MPs who took up the position against those who did not. Figure 3 has fourteen panels, seven positions each shown at entry and at departure, and the accompanying claim covers all of them: “there are no significantly diverging pre-position earnings trends for any of these positions.” The counts below are for the seven entry panels.

The coverage behind those panels is highly uneven. Table 6 counts the MPs observed at each displayed event time, from the position-timing variables distributed with the paper.

Table 6: MPs behind the seven entry panels of [Weschle \(2021\)](#) Figure 3.

Position	MPs observed			Share of year taken	
	−2	−1	0	−2	−1
Minister	36	37	65	55%	57%
Minister of state	40	40	59	68%	68%
Parliamentary secretary	71	71	97	73%	73%
Frontbench team	52	56	108	48%	52%
Committee chair	47	49	72	65%	68%
Shadow cabinet	40	81	218	18%	37%
Committee membership	27	57	415	7%	14%

Notes: Event time is years relative to taking up the position, so 0 is the year the post is taken. Shares are each pre-position count divided by the count at 0.

Coverage at the deeper pre-position point ranges from 73% to 7%. In the committee-membership panel it rests on 27 MPs against 415 at entry. One sentence therefore covers

panels where nearly the same MPs are compared before and after, and panels where a small early subset is compared with a group fifteen times larger. Figure 12 attaches those counts to the minister and committee-membership panels.




Figure 12: Two panels of Weschle (2021) Figure 3, with the MPs behind each plotted point.

Notes: Mean logged real outside earnings by event time, reproduced from the author's own script and data, with his panel titles. The comparison series is his: for each MP taking the position, the annual mean among MPs who never hold it, restricted to the years that MP is observed and re-indexed to the MP's event clock. No standard errors are reported on this display and none are added here. Rows beneath each panel give the MPs who took the position and the never-holders behind the comparison mean. All seven positions reproduce. These two span the range.

Clarke (2020): composition that shifts across the pre-period. Clarke (2020) asks whether joining an ideological faction shifts a U.S. House member's donor base, estimating a staggered difference-in-differences with candidate and congress fixed effects on Democrats between the 97th and 113th Congresses. All five columns of the paper's appendix Table A3 reproduce to the two decimals it reports, from 0.09 with no controls on 3,772 observations to 0.08 with the full set. The panel is unbalanced in the way Section 4 describes and then some: 3,772 of 12,155 possible candidate-congress cells are observed, the median candidate appears in 5 of 17 congresses, and candidates enter across 17 distinct congresses while 73% leave before the last. Only 3% have interior gaps, so the imbalance is differential entry and

exit rather than intermittent missingness.

The paper reports no event study. Figure 13 supplies one, in the author’s own specification, with the counts behind each coefficient. The treated arm runs 12, 16, 27 and 42 across the pre-period against 45 at onset, while the control arm sits between 269 and 509 throughout. The periods that would speak to parallel trends and the period that produces the estimate therefore describe different sets of legislators, and the gap is on the treated side.




Figure 13: Event study based on the data and specification of [Clarke \(2020\)](#), with the counts behind each coefficient.

Notes: Event-study interactions on Blue Dog Coalition membership with candidate and congress fixed effects, the four controls of the paper’s own specification, and standard errors clustered by candidate, with the congress before joining omitted. Event time is congresses since joining, and the tails are binned at -4 and $+4$. The 38 candidates who are already members when first observed have no pre-period and are excluded, leaving 45 with an in-sample onset. Rows beneath the panel give the treated and control candidates behind each coefficient.

The Weschle figure reuses the author’s own construction of the comparison series and his position-timing variables. The Clarke event study uses his estimation sample and specification, restricted as he restricts it to a faction’s own copartisans.

C Target Populations and Inference in the Liao (2023) Diagnostics

This appendix reports three things Section 6 states without elaborating: the rest of the parallel-trends evidence Liao (2023) offers, what the matching of Section 6.3 achieves on balance, and how many firms the fixed-composition requirement costs at each window length.

The paper’s other robustness checks. Beyond the raw-means figure, Liao (2023) reports three further checks, none of them a parallel-trends display. These are an event study presented as a column of coefficients in Appendix Table D.14, two placebo specifications in neighbouring columns of the same table, and a matched difference-in-differences for 2016–2017 using the method of Imai et al. (2023). The event-study column is the subject of Section 6.1, which reproduces it exactly and shows that its coefficients are computed from a composition that turns over almost completely from year to year. We take no position on the two placebo specifications, which ask whether the estimated effect also appears at a false onset date or for firms that lobbied in 2016. The pre-treatment path is not the relevant evidence for either.

The matched design is a two-group, two-period comparison. All fourteen of its specifications use a single pre-treatment period, so it contrasts 2016 with 2017 and nothing earlier. Matching does hold the compared population fixed, at the cost of the attrition Table 8 quantifies, and it rests on 70 treated firms. But one pre-treatment period leaves no trajectory, so the balance reported is on levels rather than trends. The paper’s evidence on *pre-trends* therefore rests on the pooled display and the event-study column, and the tables below support the discussion of both.

Balance, and the choice of control group. Table 7 reports what the matching in row D of Table 4 achieves, against the unmatched control pool and against an alternative that drops the entry-cohort constraint.

The covariates are not ours. [Liao \(2023\)](#)’s own matching code exact-matches on two-digit NAICS industry and firm size class, and assesses balance on the denial rate, industry, size, public listing, sales and employment.¹⁸ We add entry cohort, for the reason Section 4 gives, and pre-period denial rates and petition volume.

The row that matters most is the withheld one. Matching on 2011–2015 does partly flatten a 2011–2016 pre-treatment path by construction, which is the obvious objection to reading the matched event study as validation. But balance at 2016 improves too, from 0.378 to 0.134, and the score never saw that year.

Petition volume remains at a standardised difference of about 0.33 in both matched specifications, the one covariate that does not reach conventional balance, because requiring an exact cohort match while also matching on industry and size leaves a thin donor pool. The alternative specification, which handles entry cohort through the score alone, achieves marginally better balance on listing and on the held-out year while leaving a residual cohort-share difference of 0.27. We prefer the exact-matched version on the a priori ground that the control group should represent the treated group’s entry composition, not because of the estimate it produces.

What fixed composition costs, by window length. Table 8 reports how many firms survive the fixed-composition requirement at each window length. Each row requires a firm to be observed in every year of a window running back from 2017, and counts what survives against the 74 treated and 38,188 control firms observed at onset.

Treated retention declines gently, from 88% to 64%, as the window lengthens to sixteen pre-treatment years, then falls away sharply to 20% at the full published span. The reason

¹⁸The exact-matching variables are two-digit NAICS industry and firm size class, taken from the author’s own matching code. The paired specifications with and without that constraint are columns 1 and 2 of his Appendix Table D.9, which for our treatment definition give -2.9 and -1.71 percentage points. We use industry, size and listing, recorded for 93% of the control firms and 98% of the treated firms under fixed composition. Sales and employment are populated too sparsely to match on. These firm characteristics are not distributed with the author’s replication package. They are our own extract, retrieved from Orbis on 23–24 August 2026 by matching on Bureau van Dijk identifiers, and are not necessarily the vintage used in the published analysis.

Table 7: What the matching achieves. Standardised differences and composition gaps between treated and control firms, on the fixed-composition sample.

	Full control pool	Propensity score matched	
		Plus exact on cohort	Cohort in score only
Control firms	5,854	128	128
<i>Composition: total absolute difference in category shares</i>			
Entry cohort	1.00	0.00	0.27
Industry	1.02	0.22	0.23
<i>Absolute standardised difference in means</i>			
Denial rate, 2011–15	0.515	0.019	0.169
Denial rate, 2016 (held out)	0.378	0.134	0.125
Petition volume	1.151	0.325	0.330
Size class	1.532	0.039	0.000
Listed	1.577	0.287	0.144

Notes: The 64 Orbis-merged treated firms under fixed composition. Column 2 is row D of Table 4. The composition measure runs from 0 when the treated and control distributions coincide to 2 when they are disjoint. It is not a standardised difference, since neither variable has a scalar mean. Matching is nearest-neighbor, two control firms per treated firm without replacement, on a propensity score over every row shown except the bold one. Both matched columns put entry cohort in the score, and the second also matches on it exactly. The bold row is therefore an out-of-sample check.

Table 8: Treated and control firms retained, by the length of the window over which continuous observation is required.

Window	Pre-periods	Treated	Control	Treated retained
2012–2017	5	65	7,294	88%
2011–2017	6	65	6,303	88%
2009–2017	8	60	4,651	81%
2007–2017	10	52	3,672	70%
2005–2017	12	50	2,768	68%
2003–2017	14	49	1,981	66%
2001–2017	16	47	1,497	64%
1999–2017	18	35	707	47%
1997–2017	20	25	332	34%
1995–2017	22	21	179	28%
1992–2017	25	15	49	20%

is structural, as a window beginning in 2001 can accommodate any firm first observed in or before 2001, whereas one beginning in 1992 can only contain firms first observed in 1992, so the deep windows are limited by entry timing rather than by attrition. The control group falls much faster than the treated group throughout, which is the same asymmetry the published display exhibits. We report the whole range because it removes the need to defend a particular window.

D Additional Empirical-Application Diagnostics

This appendix collects supporting diagnostics for the within-cohort matching estimator of Section 7.3: the candidate-level pre-treatment outcomes that inform the trimming step (Figure 14), the matched within-cohort raw trajectories (Figure 15), the cohort event studies carried one period past common exit (Figure 16), and the within-cohort covariate balance after matching (Figure 17).




Figure 14: Candidate-level pre-treatment outcomes, by entry cohort.

Notes: Each point is a legislator’s pre-treatment log corporate outcome in a given Congress, faceted by entry cohort. Color distinguishes defectors ($D114 = 1$) from non-defecting controls ($D114 = 0$). Defectors whose pre-period outcomes fall outside the support of the controls in their cohort are dropped before matching, two in the 112th cohort and one in the 113th. Each marker is a single legislator–Congress observation, so no confidence intervals are shown.




Figure 15: Within-cohort raw trajectories for matched defectors and their controls.

Notes: Mean log corporate fundraising for matched defectors and non-defectors, by Congress, faceted by entry cohort. Bars are ± 1.96 times the standard error of each group mean within a Congress. That is the sampling error of the plotted raw mean, unclustered, not the legislator-clustered interval of the event-study figures. The dashed line marks treatment onset at the 114th Congress.




Figure 16: The cohort event studies carried one period past common exit.

Notes: Figure 11 extended by one post-treatment period, with specification, reference period and matching unchanged, so the pre-treatment coefficients are identical. The Congresses after the 114th are averaged into a single observation, so every legislator still serving contributes the same post-treatment length. That equalizes duration but not presence. Four defectors leave the House after the revolt and appear in no post-114 period, so the final point in each panel rests on fewer legislators than the rest and the printed counts describe every displayed period except the last.




Figure 17: Within-cohort covariate balance after matching.

Notes: Distributions of district vote share (DVS) and ideology (dynamic CF score) for matched defectors and non-defectors, by entry cohort. Balance is close in most cohorts. The 114th cohort is the least well balanced, with three treated units and a single pre-period. Dropping it barely moves the aggregate estimate. Boxes span the interquartile range with the median line, and whiskers extend to $1.5\times$ the interquartile range. No confidence intervals are shown.